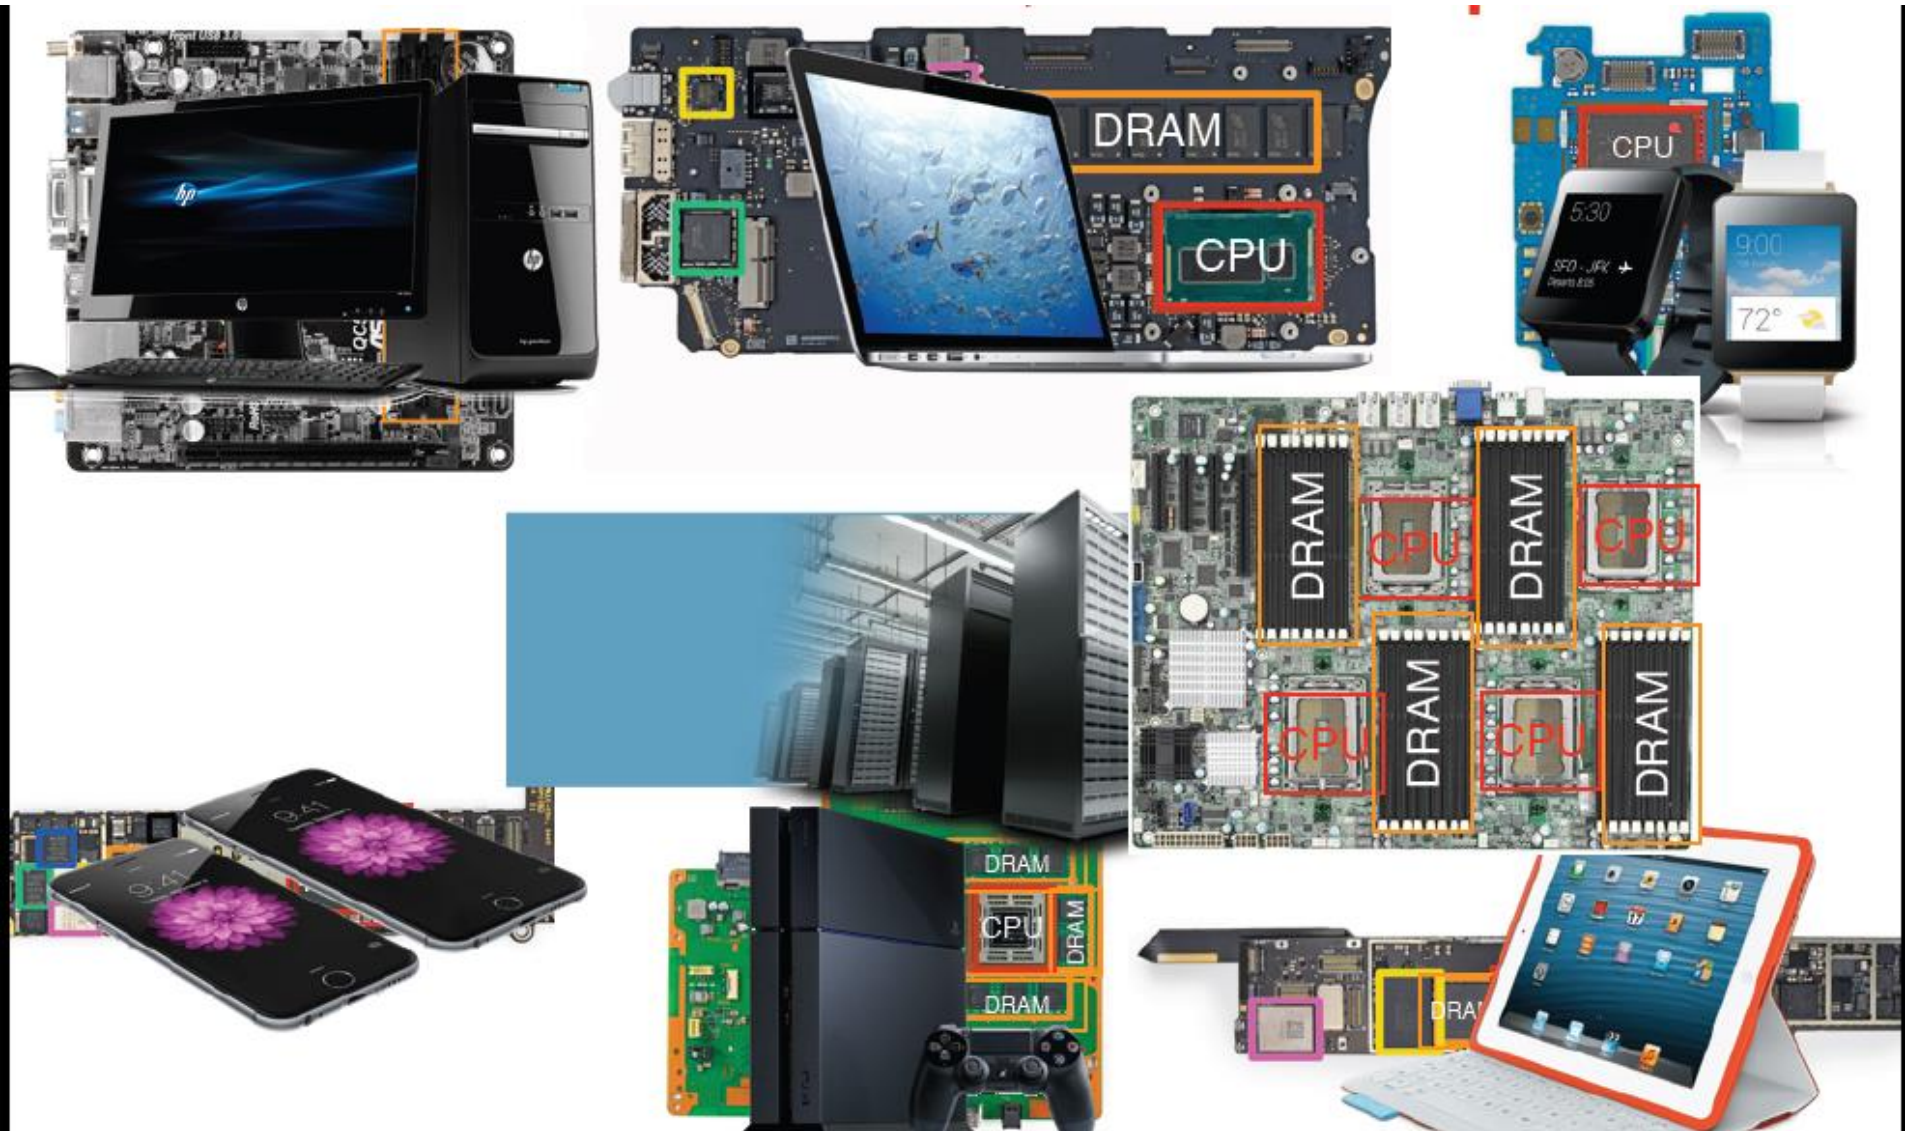
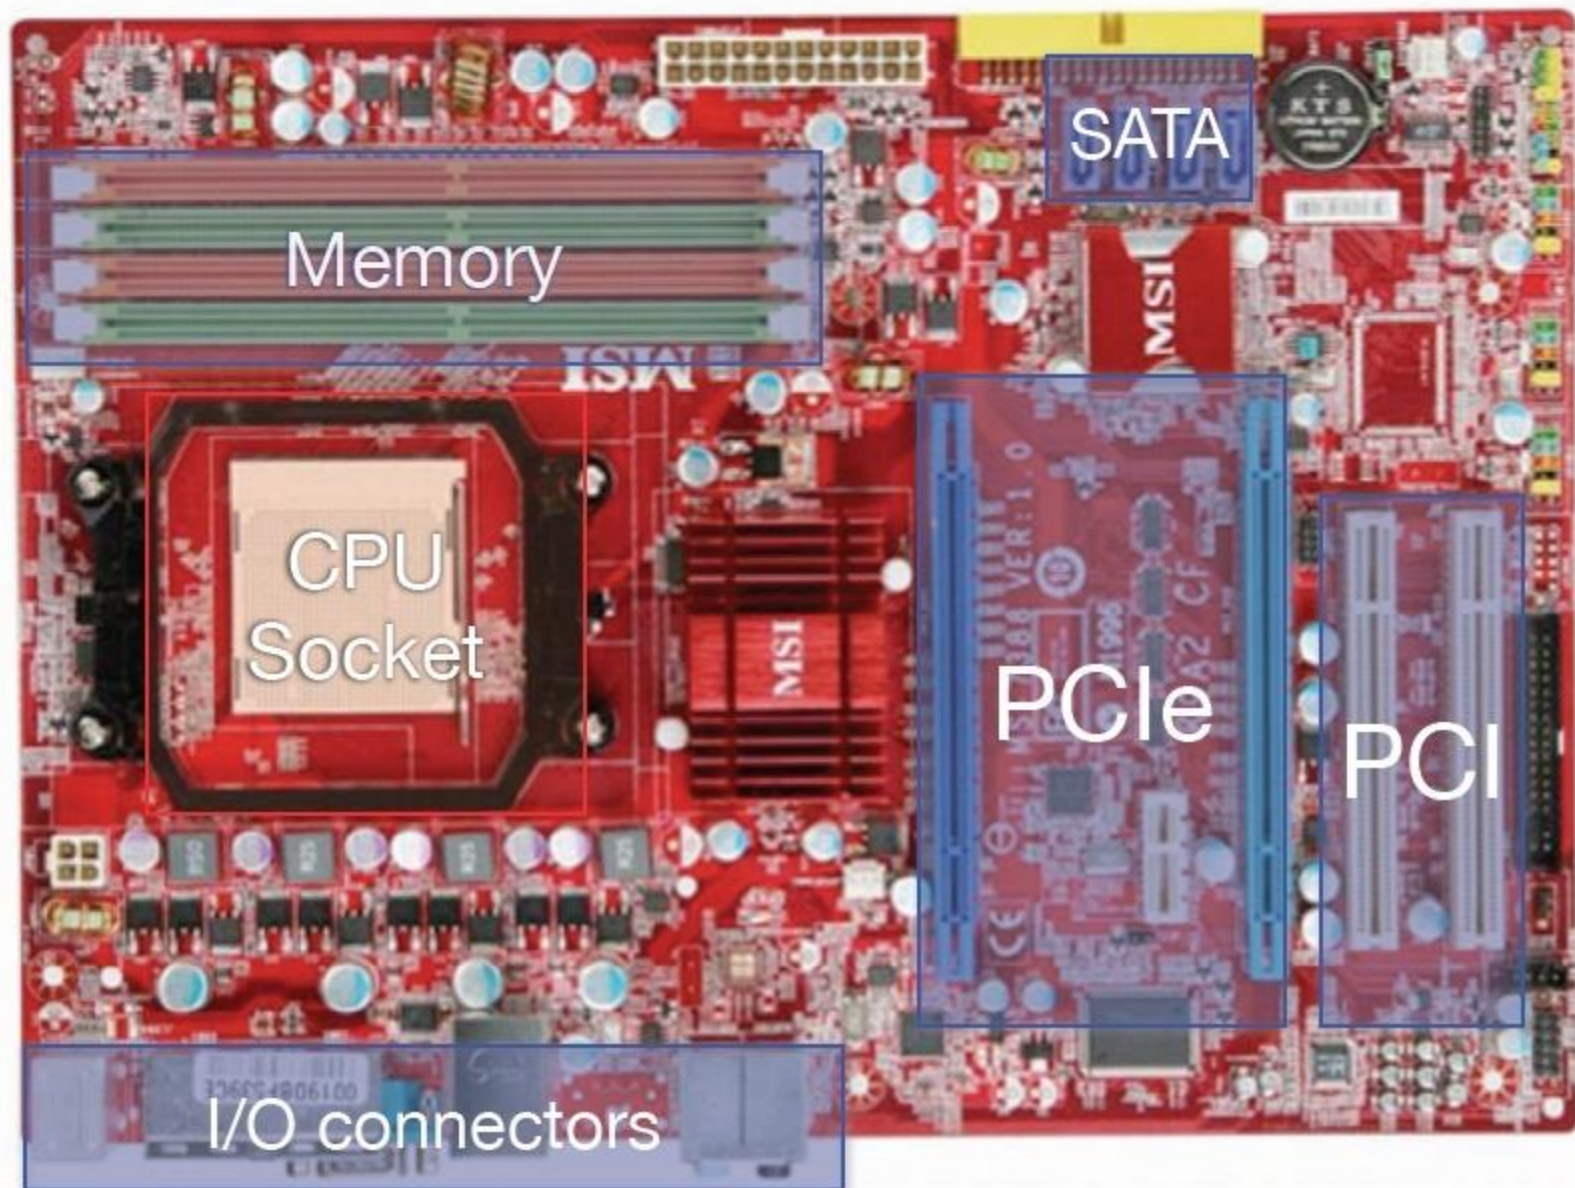


# Systems

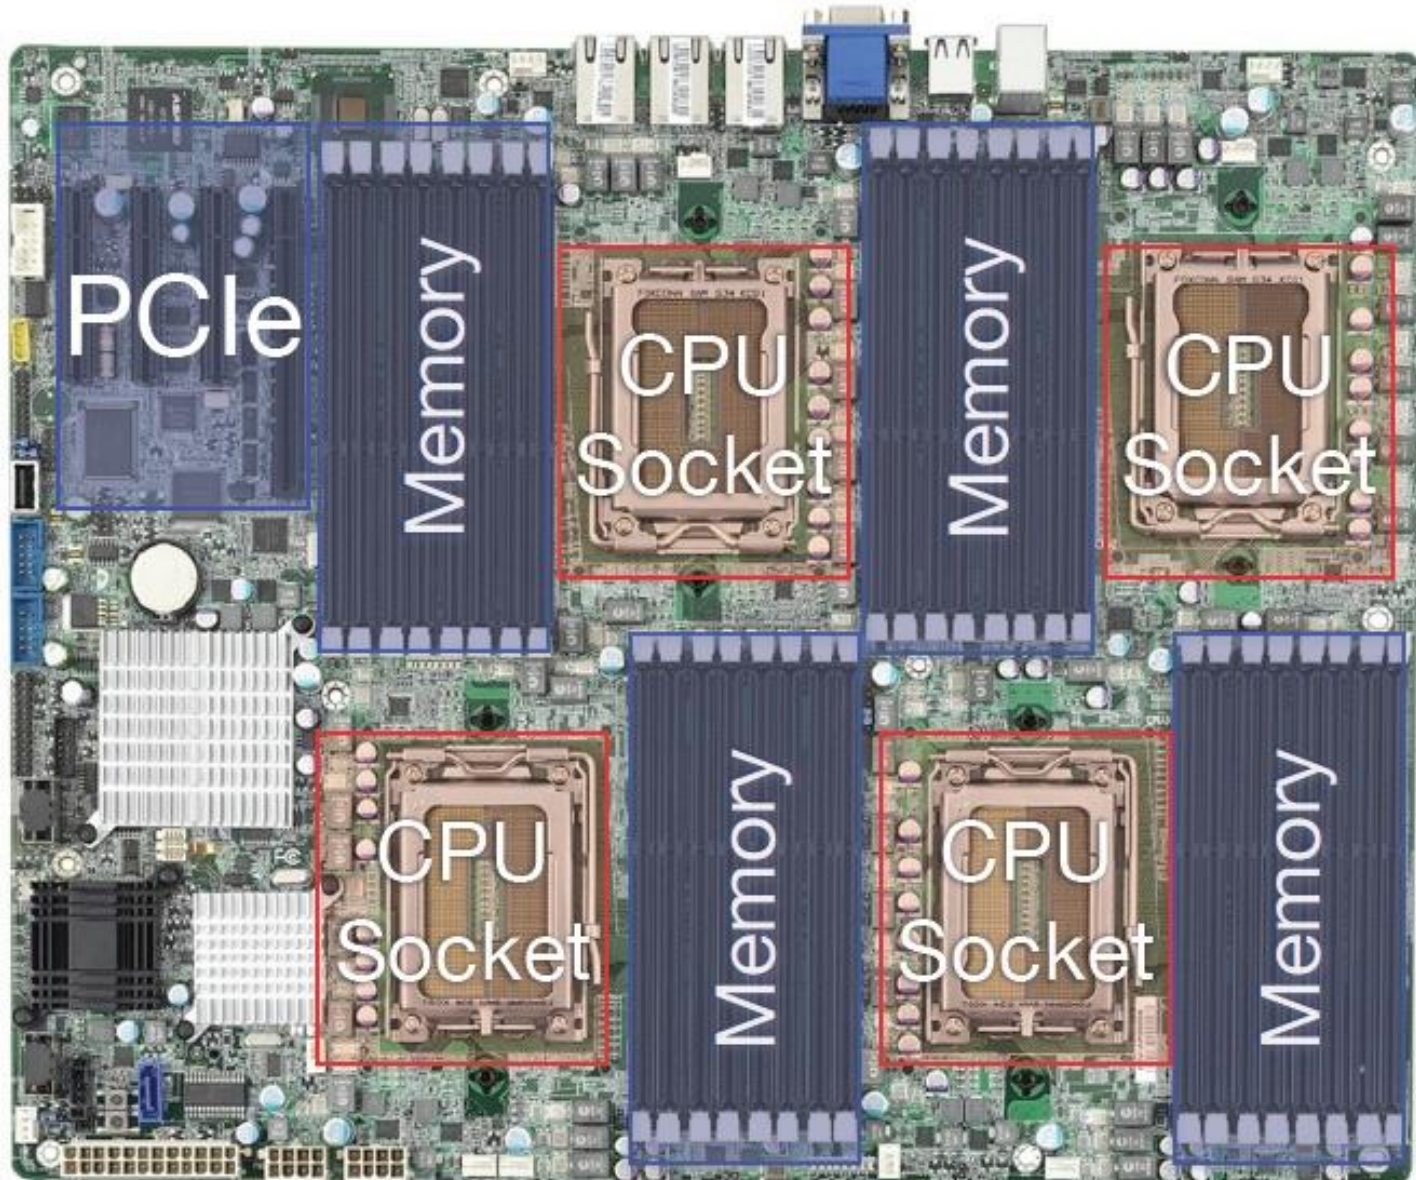


# Desktop PC

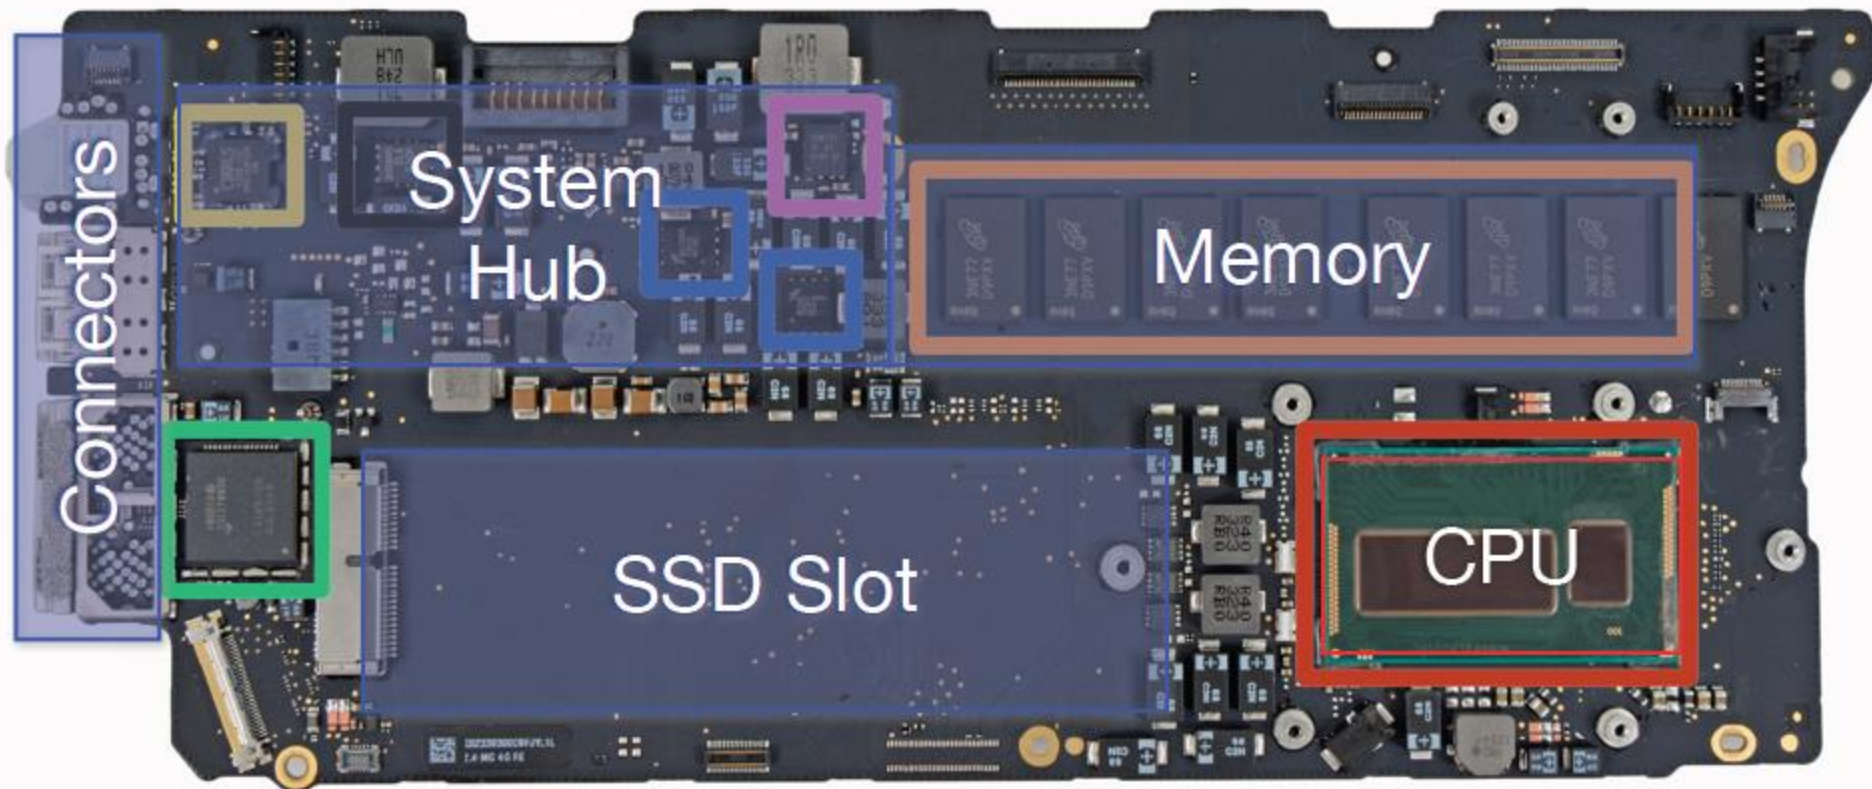




# Server

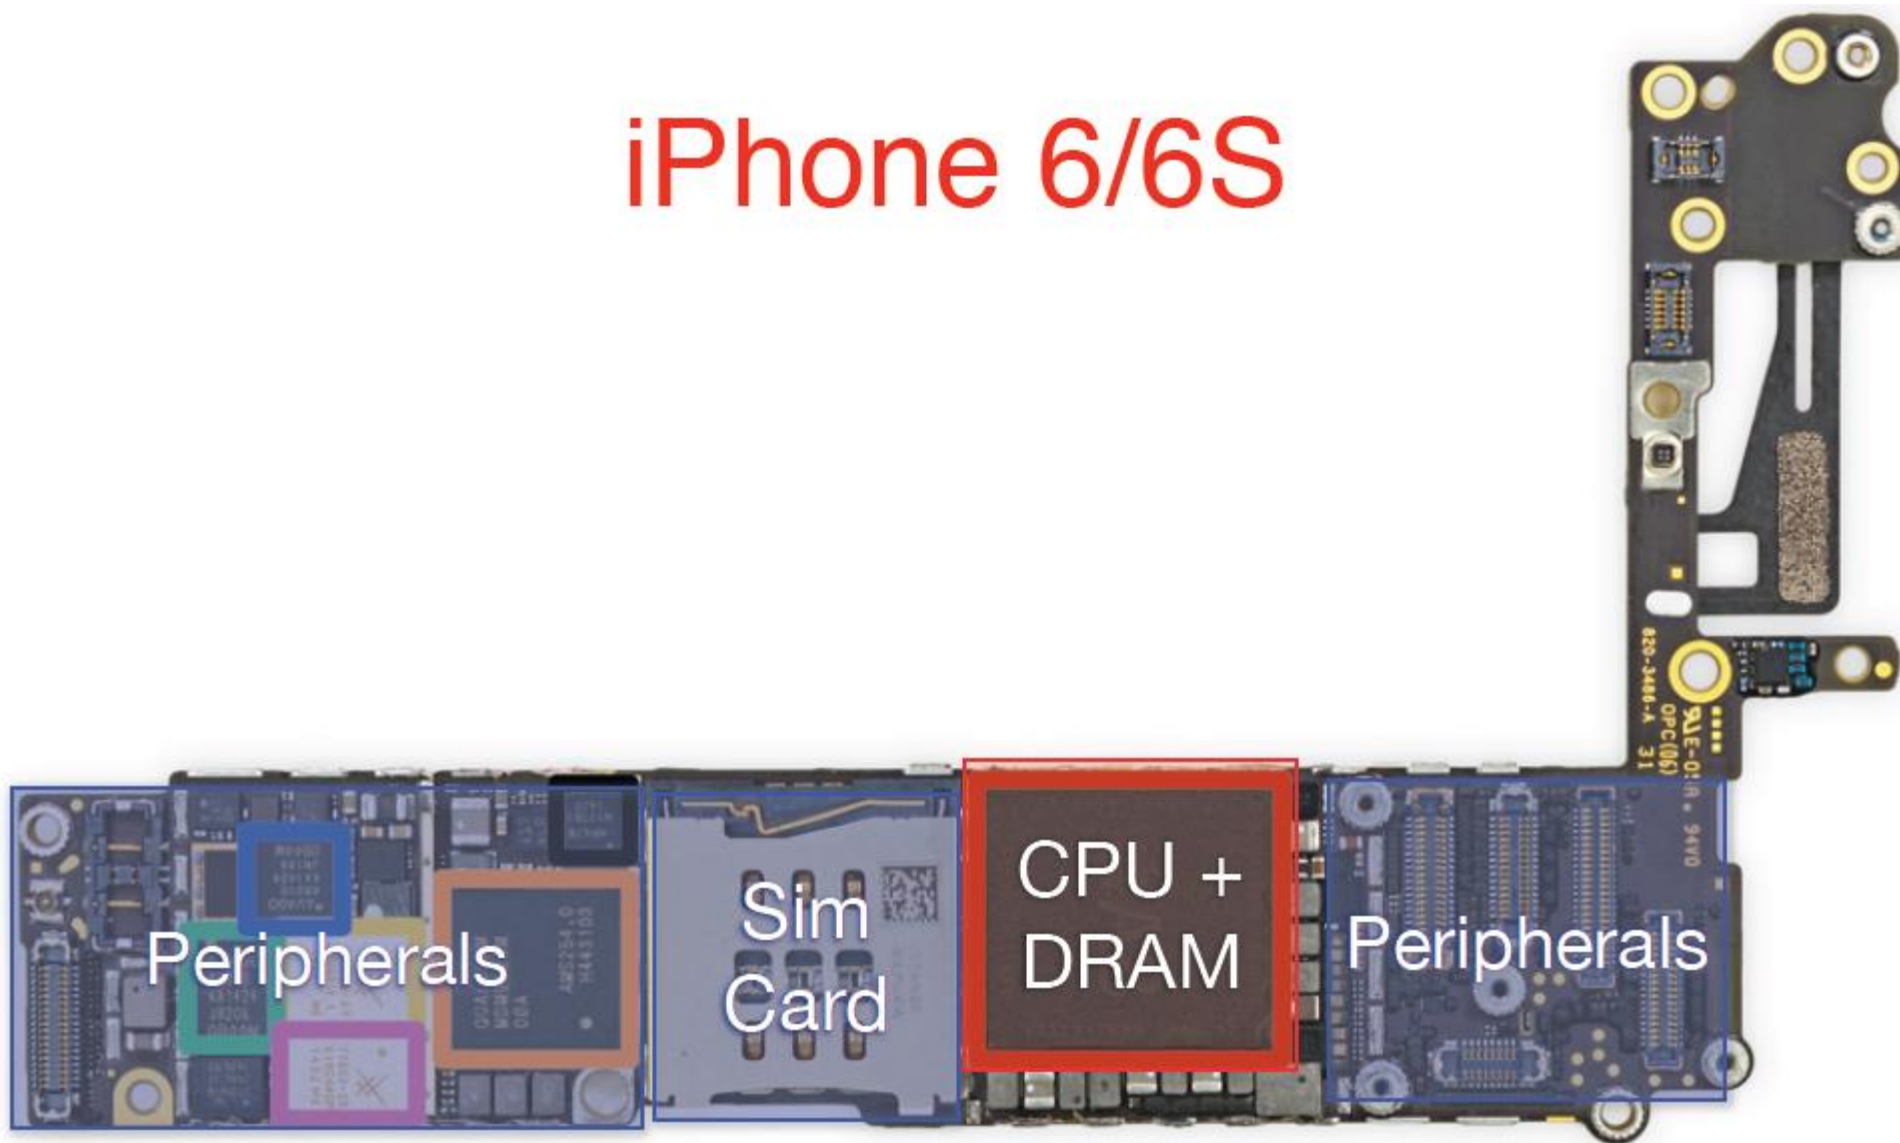


# Macbook Pro w/ Retina





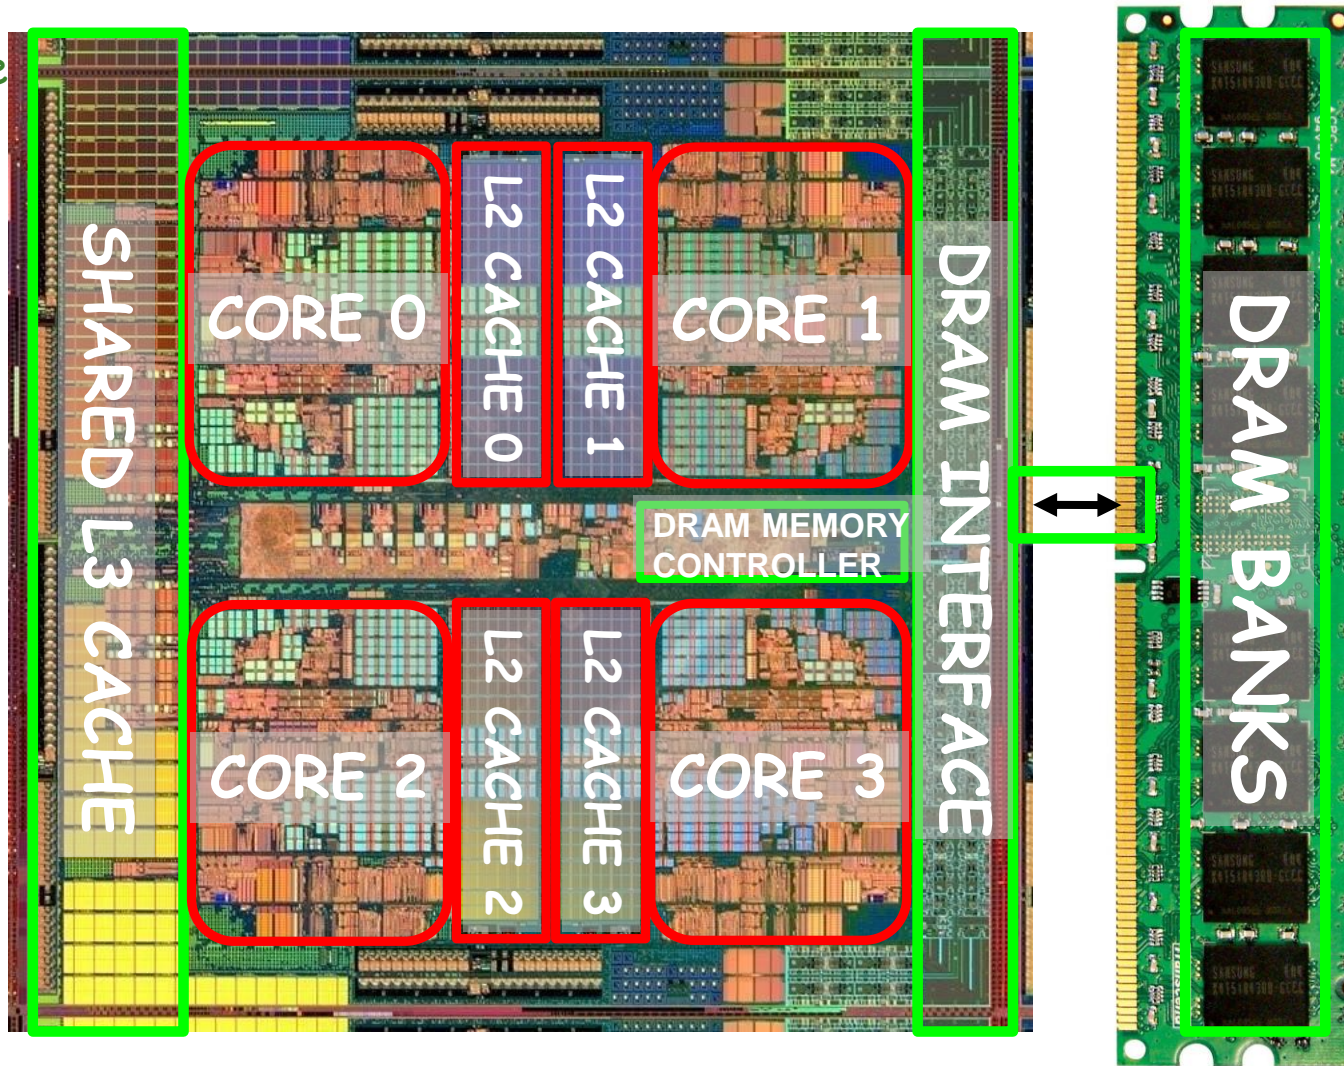
# iPhone 6/6S



# Memory in a Modern System


# An Example: Multi-Core Systems

Multi-Core  
Chip



\*Die photo credit: AMD Barcelona

# Unexpected Slowdowns in Mul



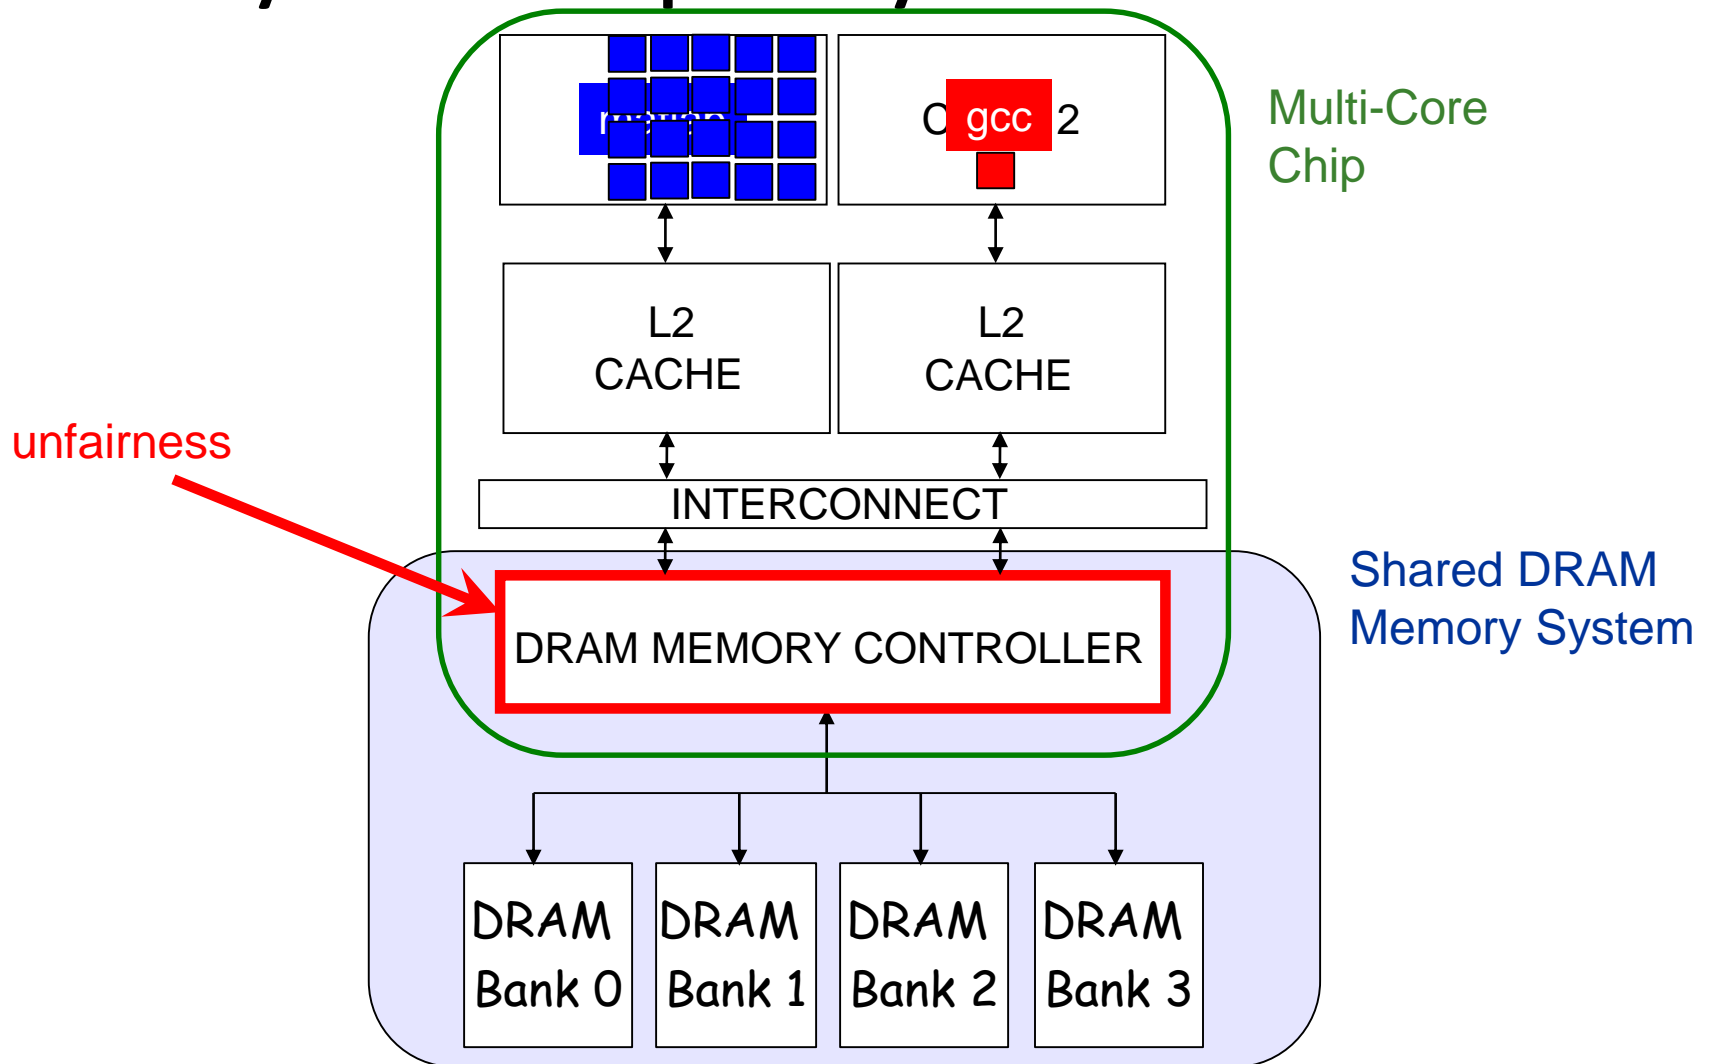
Moscibroda and Mutlu, “[Memory performance attacks: Denial of memory service in multi-core systems](#),” USENIX Security 2007.



## A Question or Two

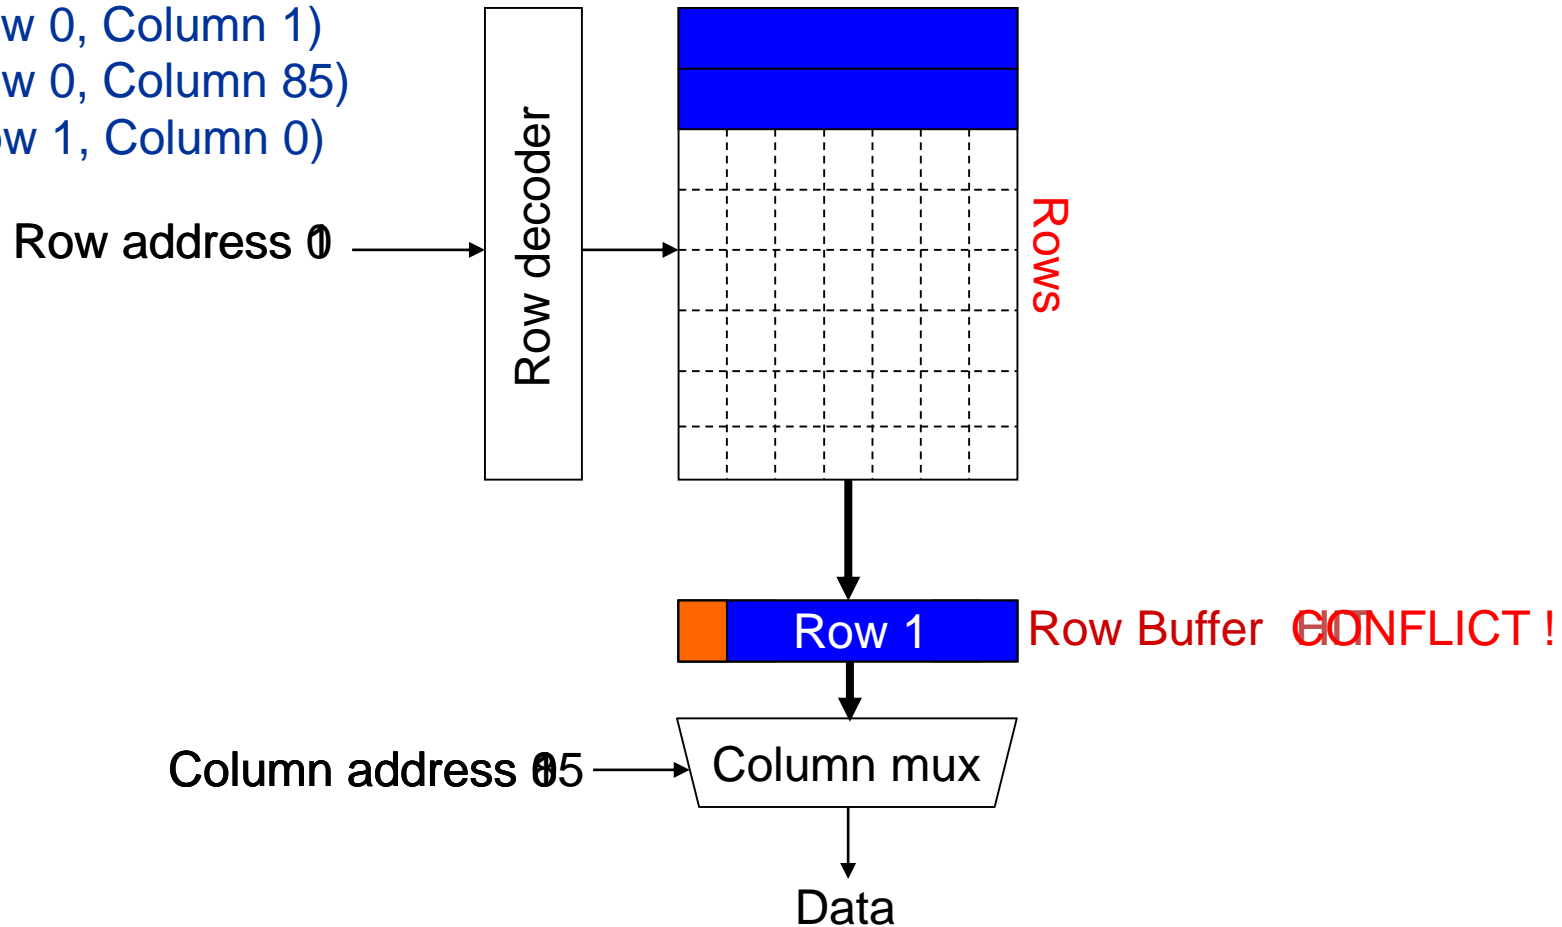
- Can you figure out why there is a disparity in slowdowns if you do not know how the processor executes the programs?
- Can you fix the problem without knowing what is happening “underneath”?

# Why the Disparity in Slowdowns?



# DRAM Bank Operation

Access Address:  
(Row 0, Column 0)  
(Row 0, Column 1)  
(Row 0, Column 85)  
(Row 1, Column 0)






# Overview

- **Review of Memory Technologies**
- **Overview of Memory Hierarchy**
- **Cache Design Principles**
- **Learning Objectives**
  - Why is that some memories slow ?
  - What is memory hierarchy?
  - Why do we need memory hierarchy?
  - What is a Cache?

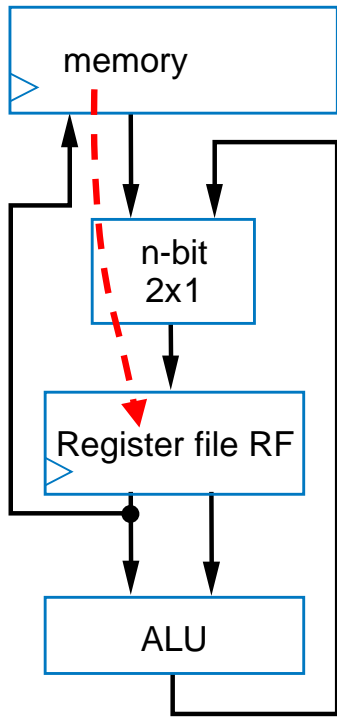
## Processor-DRAM Memory Gap (latency)



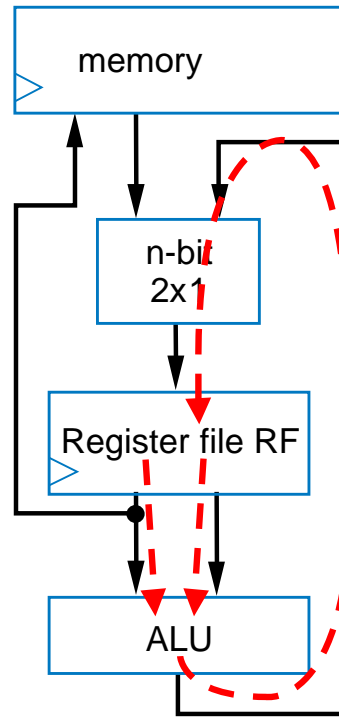
Source: David Patterson, UCB

# Review: Datapath Operations

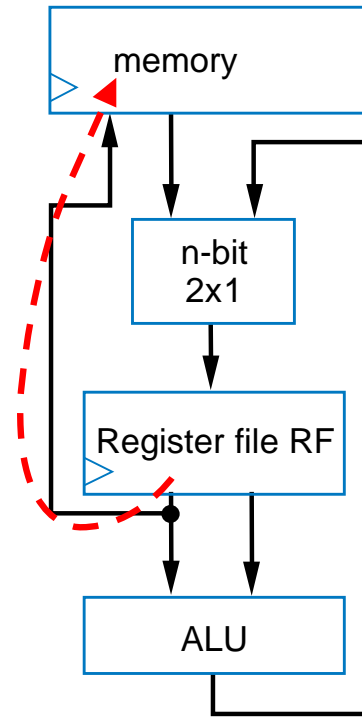
- Load operation: Load data from data memory to RF
- ALU operation: Transforms data by passing one or two RF register values through ALU, performing operation (ADD, SUB, AND, OR, etc.), and writing back into RF.
- Store operation: Stores RF register value back into data memory
- Each operation can be done in one clock cycle




Load operation



ALU operation




Store operation





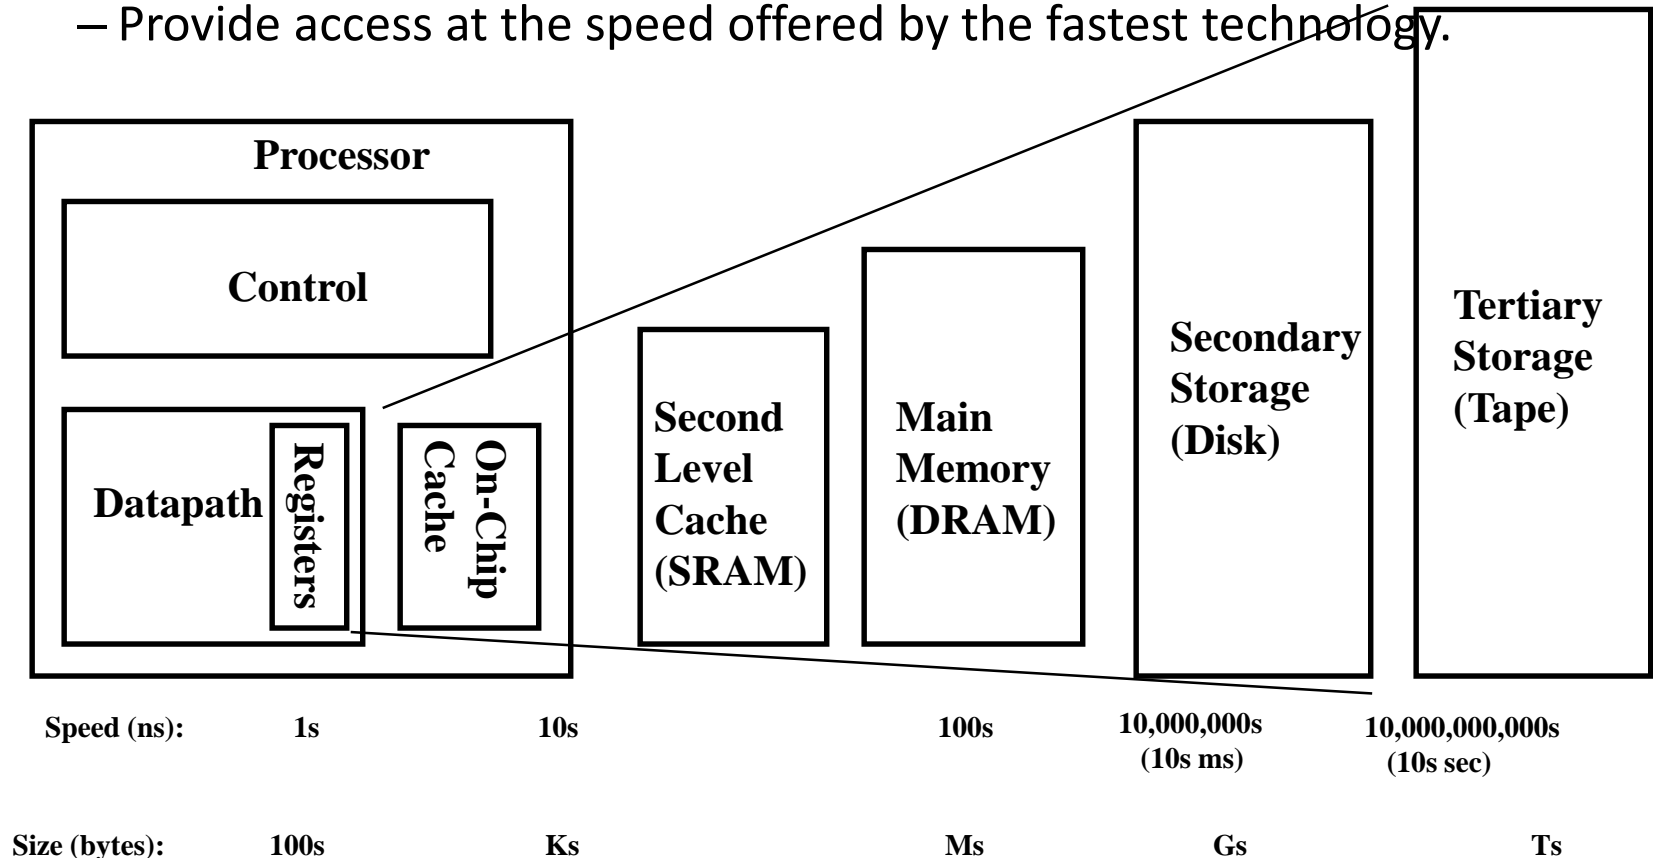
# **What are the Common Memory Technologies?**

# Memory



# Memory Hierarchy of a Modern Computer System

- By taking advantage of the principle of locality:
  - Present the user with as much memory as is available in the cheapest technology.
  - Provide access at the speed offered by the fastest technology.





# Memory

Address	Data
000000	00111110
000001	01101011
000010	01011101
000011	01100011
000100	00111110
000101	00000000
000110	11111111
000111	01010101
001000	10101010
001001	00100001
001010	11011010

64x8 RAM	
A3	D7
A2	D6
A1	D5
A0	D4
	D3
	D2
	D1
Write	D0

# 8x4 RAM

Address

Data

000



001



010



011



100



101



110



111




A2

A1

A0

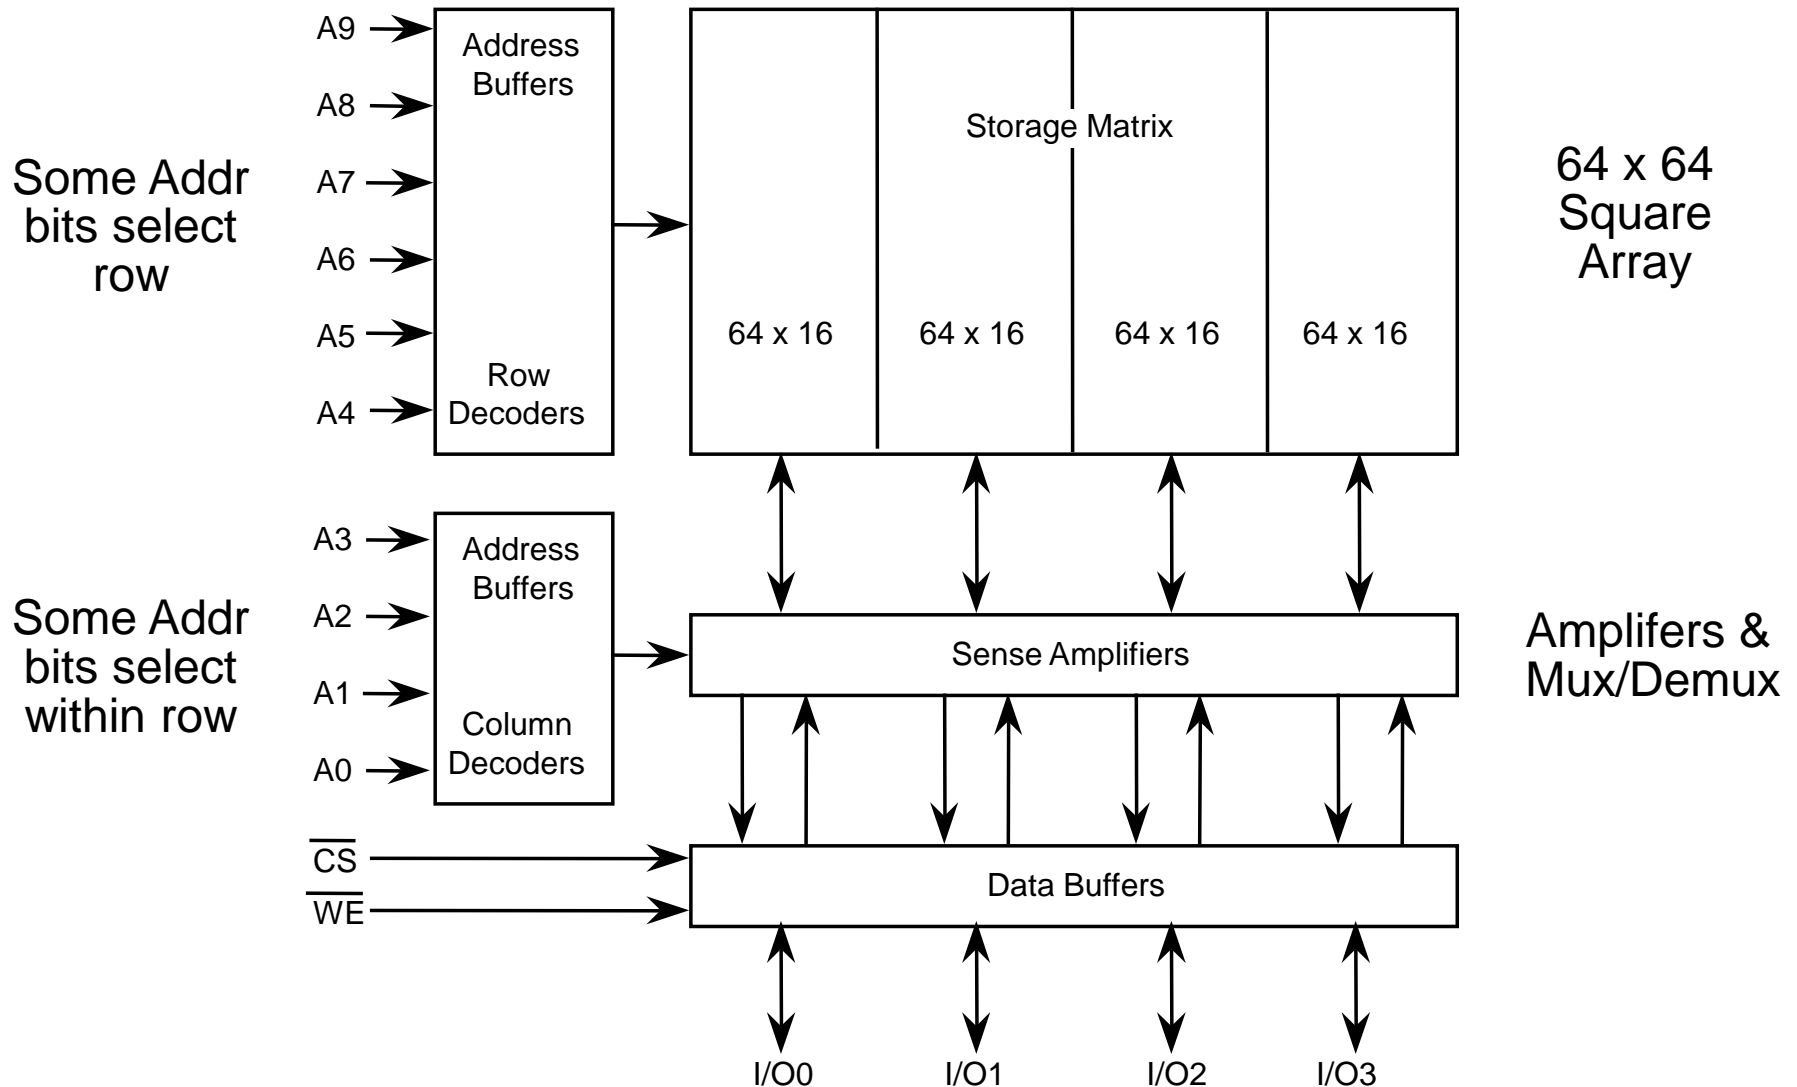
# Static RAM Organization

Chip Select Line (active lo)  
Write Enable Line (active lo)  
10 Address Lines  
4 Bidirectional Data Lines



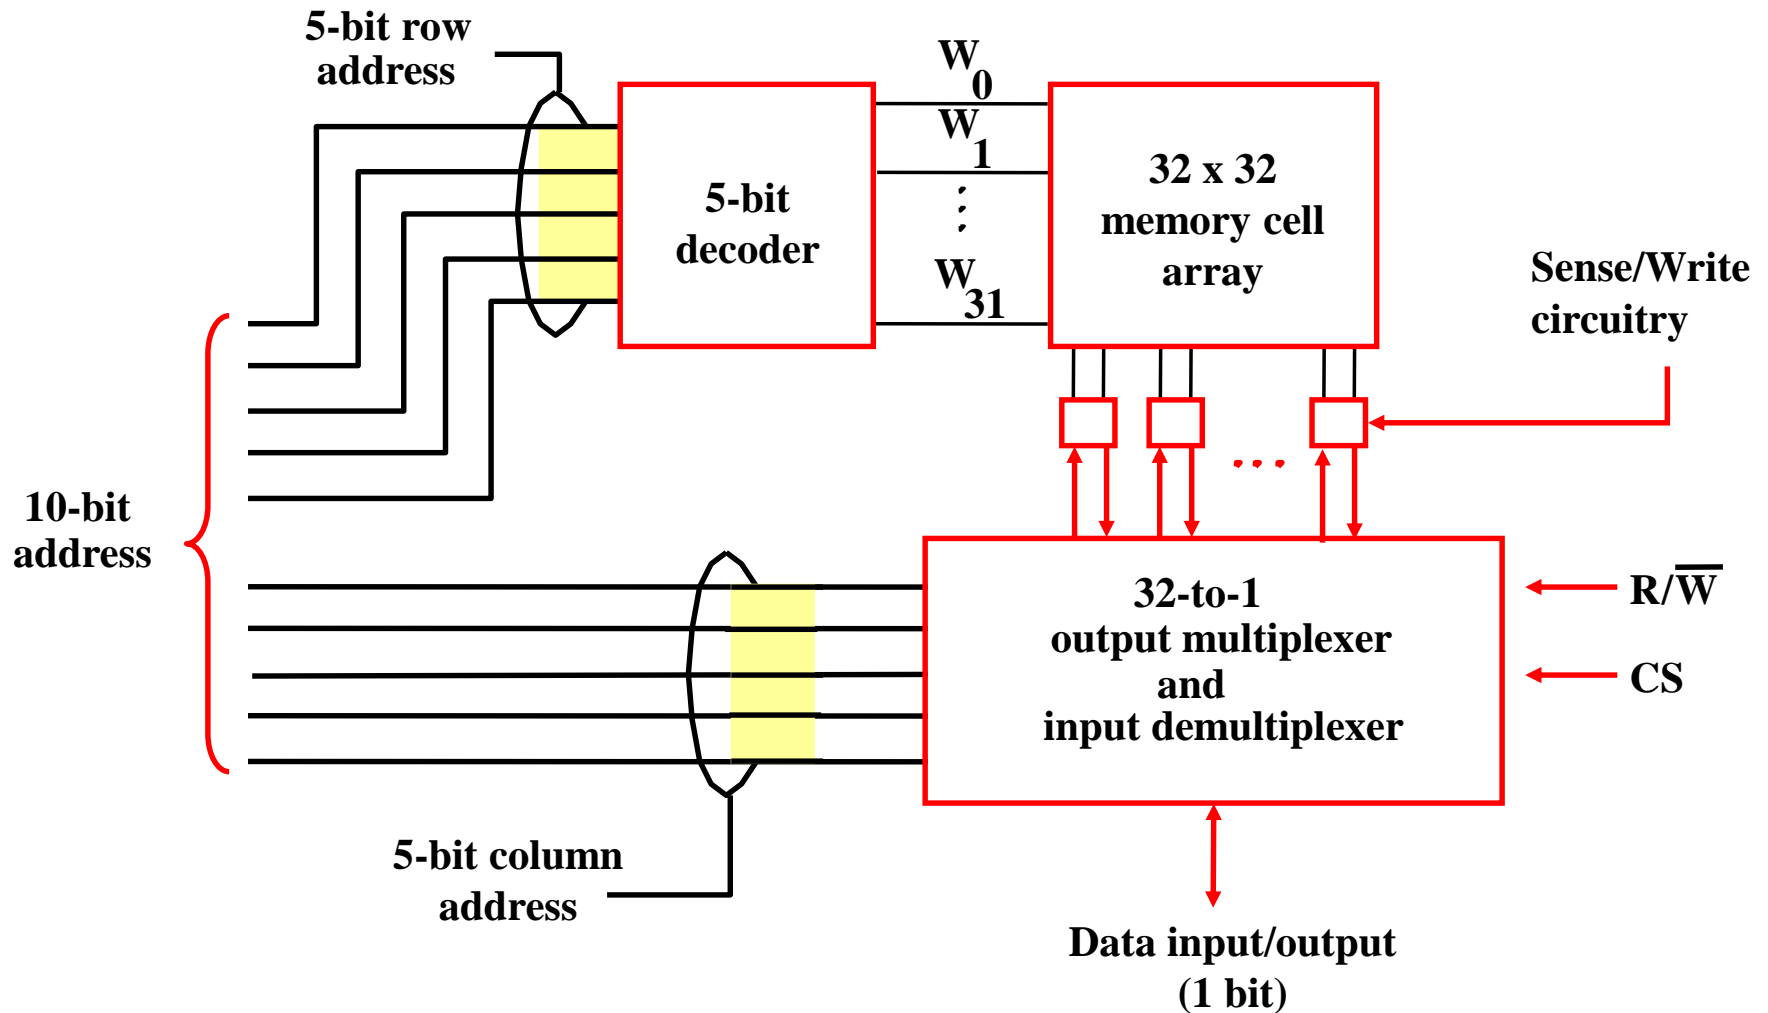
# RAM Organization

Long thin layouts are not the best organization for a RAM





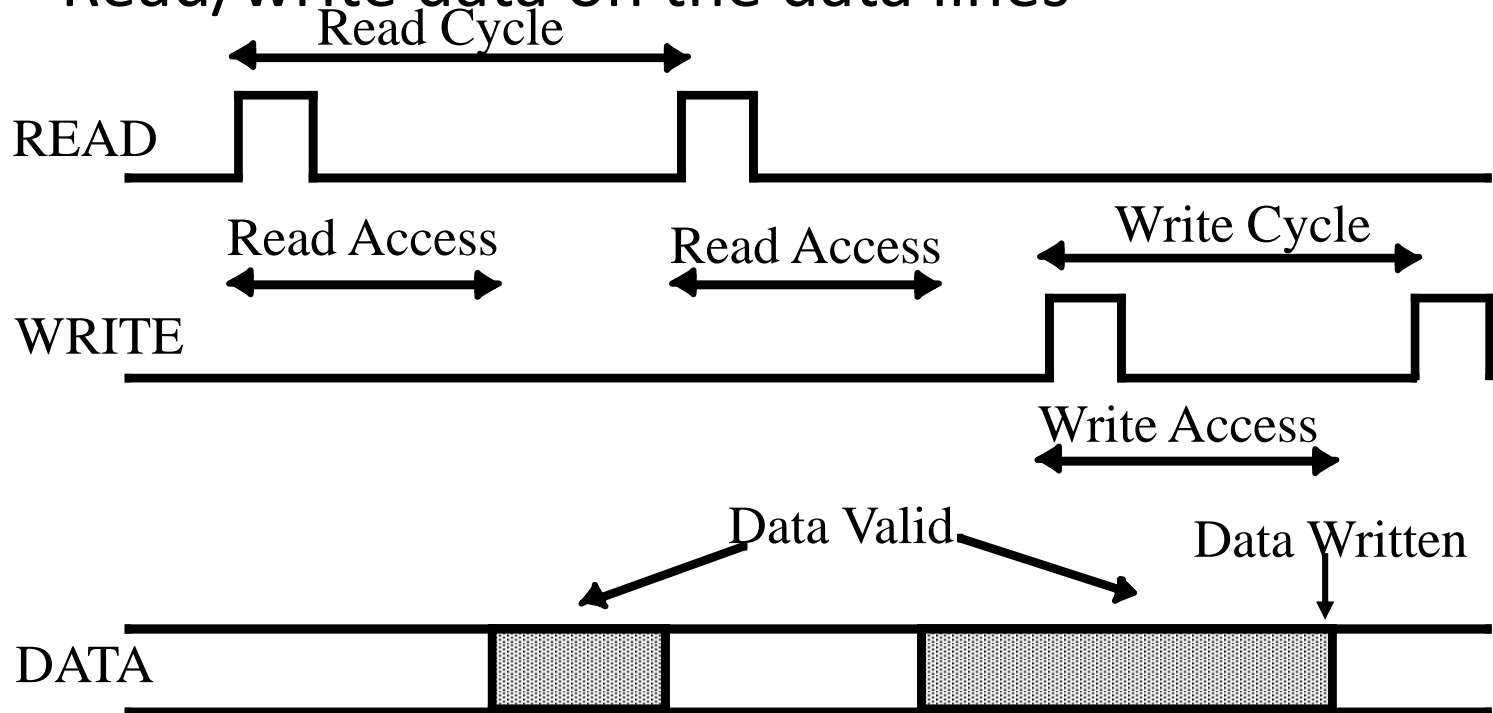
# 1Kx 1 bit RAM Organization



Organization of a  $1K \times 1$  memory chip.

# Memory Access Timing: the Big Picture

- Timing:
  - Send address on the address lines, wait for the word line to become stable
  - Read/write data on the data lines



# Content of a memory

- Each **word in memory** is assigned an **identification number**, called an **address**, starting **from 0 up to  $2^k-1$** , where **k** is the **number of address lines**.
- The number of words in a memory with one of the letters  $K=2^{10}$ ,  $M=2^{20}$ , or  $G=2^{30}$ .

$$64K = 2^{16} \quad 2M = 2^{21}$$


$$4G = 2^{32}$$

Memory address		Memory content
Binary	decimal	
0000000000	0	1011010101011101
0000000001	1	1010101110001001
0000000010	2	0000110101000110
	⋮	⋮
1111111101	1021	1001110100010100
1111111110	1022	0000110100011110
1111111111	1023	1101111000100101

1024x16 Memory Module

# Memory Cell Array Access Example

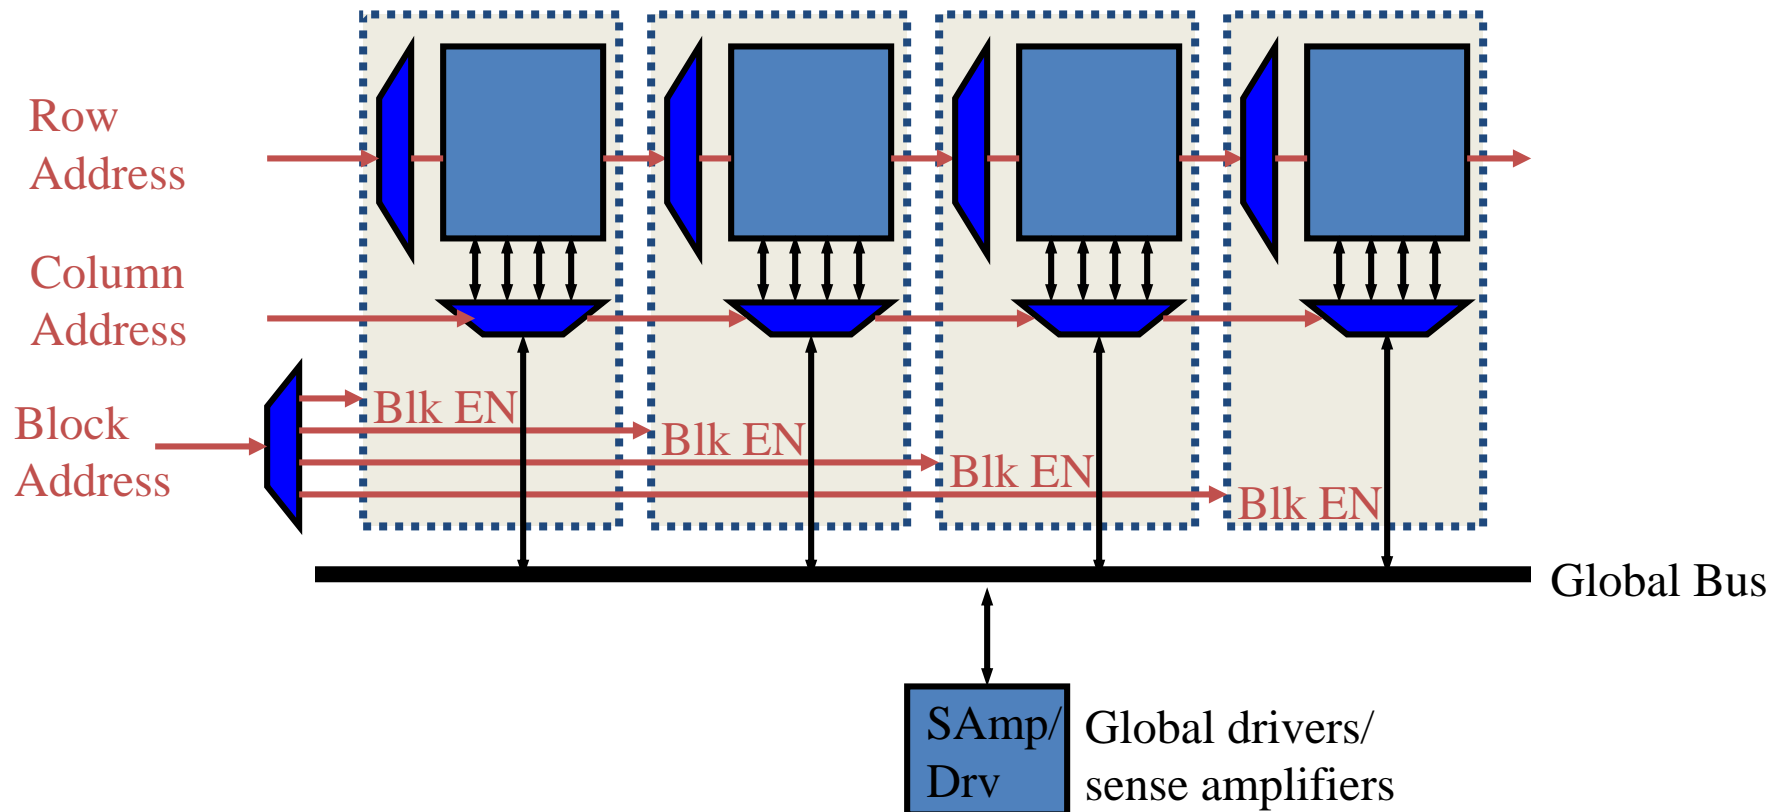
- word=16-bit wide(N), row=8 words(L), address=10 bits (k)
- Accessing word 9=  $0000001001_2$






# Memory Structures

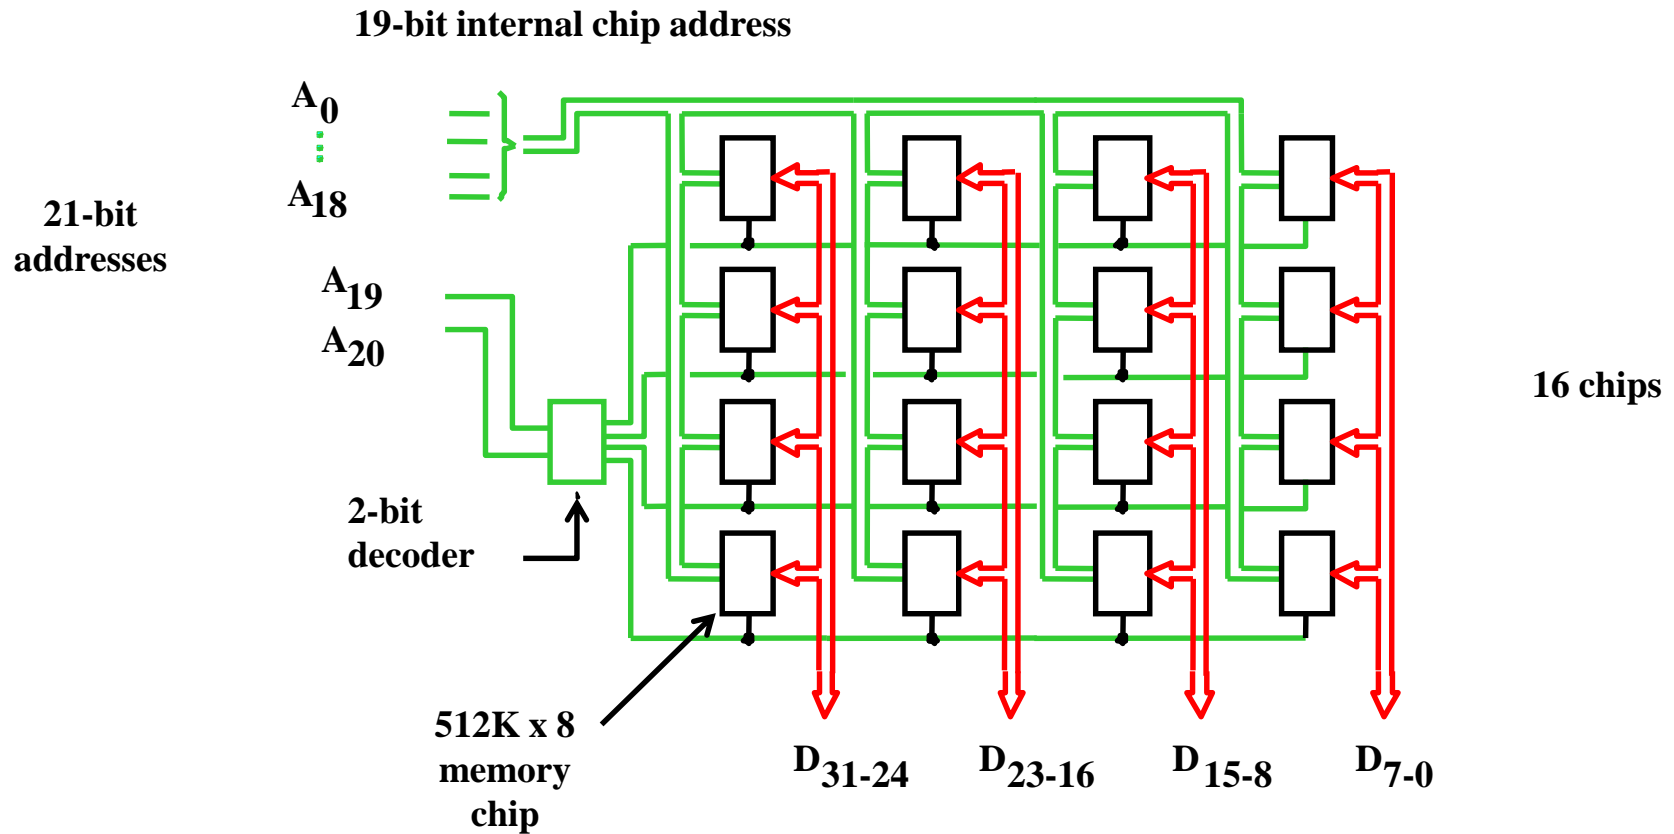
- Taking the idea one step further
  - Shorter wires within each block
  - Enable only one block addr decoder → power savings



# Larger Memories Using Multiple Chips




Question ? Design a 2Mx32 given: 512kx8



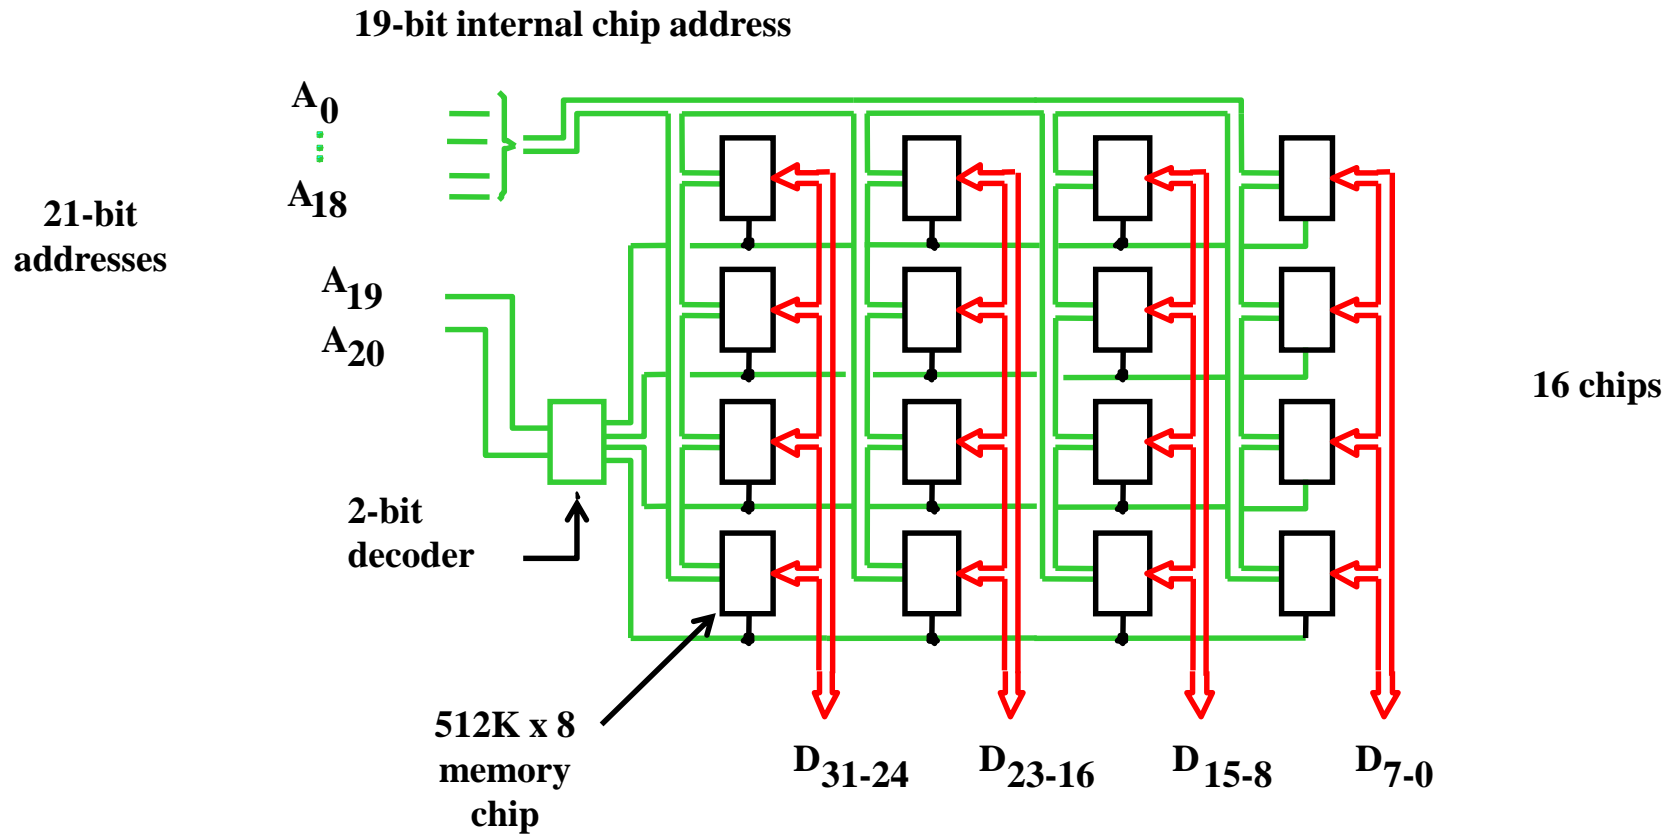
Organization of a  $2M \times 32$  memory module using  $512K \times 8$  static memory chips (16 chips).

## 19-bit internal chip address



Organization of a  $2M \times 32$  memory module using  $512K \times 8$  static memory chips (16 chips).







Organization of a  $2M \times 32$  memory module using  $512K \times 8$  static memory chips (16 chips).

# Random Access Memory (RAM)

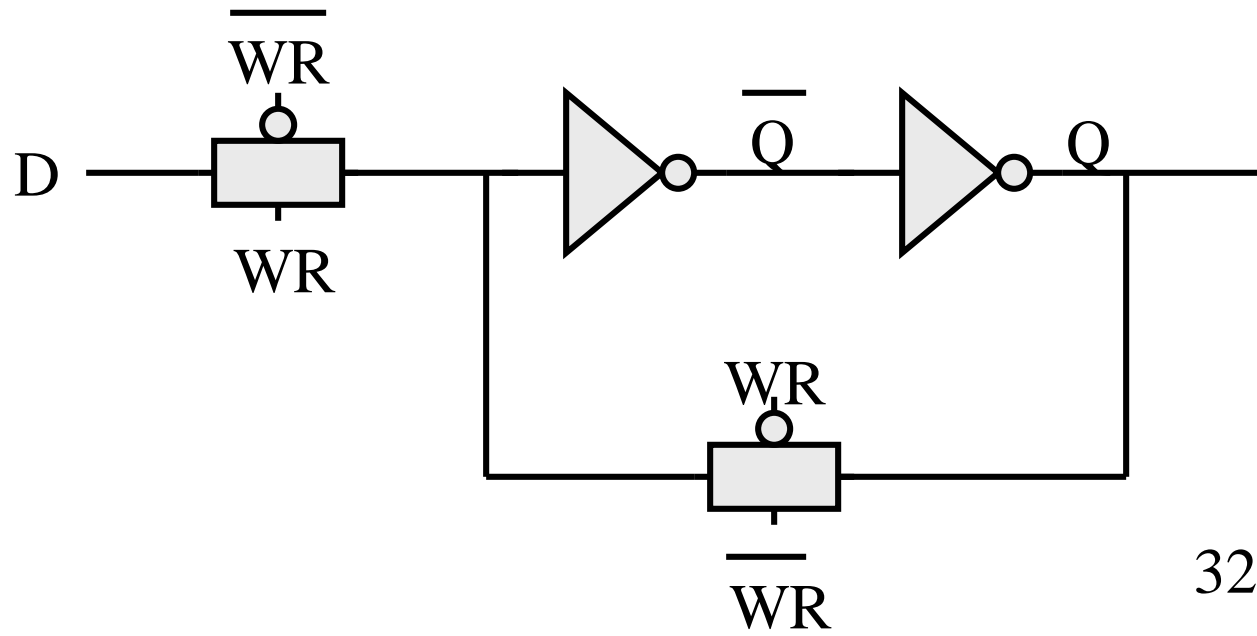
- RAM – Readable and writable memory
  - “Random access memory”
    - Strange name—Created several decades ago to contrast with sequentially-accessed storage like tape drives
  - Logically same as register file—Memory with address inputs, data inputs/outputs, and control
    - RAM usually one port; RF usually two or more
  - RAM vs. RF
    - RAM typically larger than *about* 512 or 1024 words
    - RAM typically stores bits using a bit storage approach that is more efficient than a flip-flop
    - RAM typically implemented on a chip in a square rather than rectangular shape—keeps longest wires (hence delay) short




RAM block symbol

# Implementing Registers in CMOS

- Uses transmission gate
  - When “WR” asserted, “write” operation will take place
  - Stack D latch structures to get n-bit register



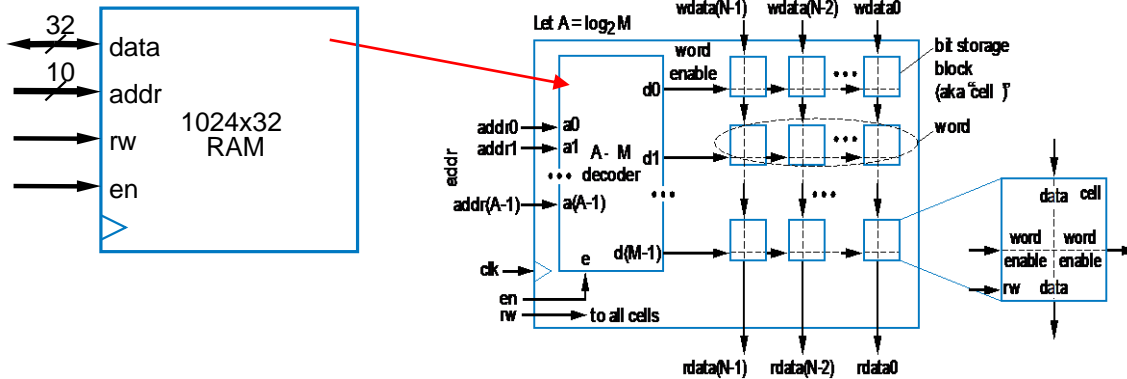
# RAM Internal Structure



Similar internal structure as register file


- Decoder enables appropriate word based on address inputs
- $rw$  controls whether cell is written or read
- $rd$  and  $wr$  data lines typically combined
- Let's see what's inside each RAM cell

# Static RAM (SRAM)





- “Static” RAM cell
  - 6 transistors (recall inverter is 2 transistors)
  - Writing this cell
    - *word enable* input comes from decoder
    - When 0, value *d* loops around inverters
      - That loop is where a bit stays stored
    - When 1, the *data* bit value enters the loop
      - *data* is the bit to be stored in this cell
      - *data'* enters on other side
      - Example shows a “1” being written into cell

SRAM cell

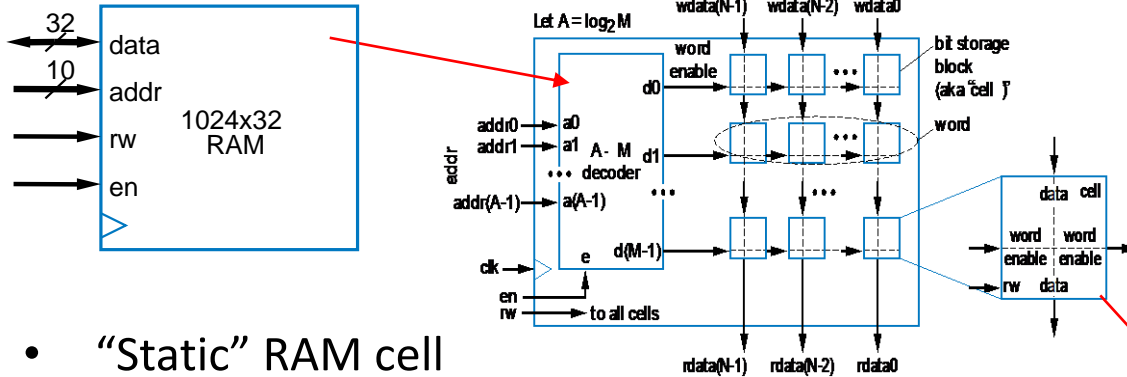


SRAM cell



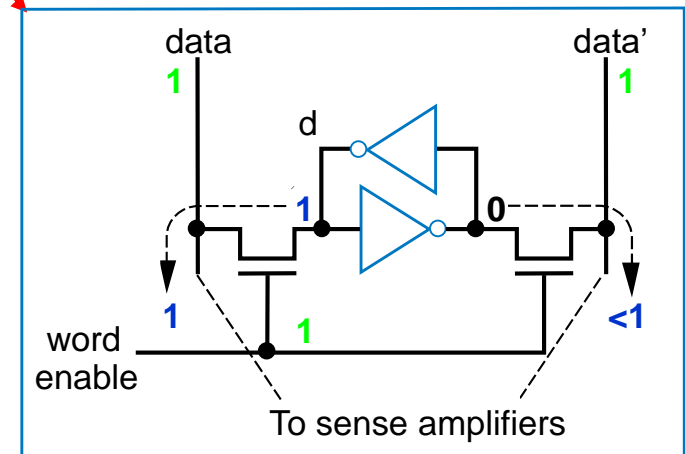


# Static RAM (SRAM)




- “Static” RAM cell
  - Reading this cell
    - Somewhat trickier
    - When rw set to read, the RAM logic sets both *data* and *data'* to 1
    - The stored bit *d* will pull either the left line or the right bit down slightly below 1
    - “Sense amplifiers” detect which side is slightly pulled down
  - The electrical description of SRAM is really beyond our scope – just general idea here, mainly to contrast with *DRAM*...

SRAM cell

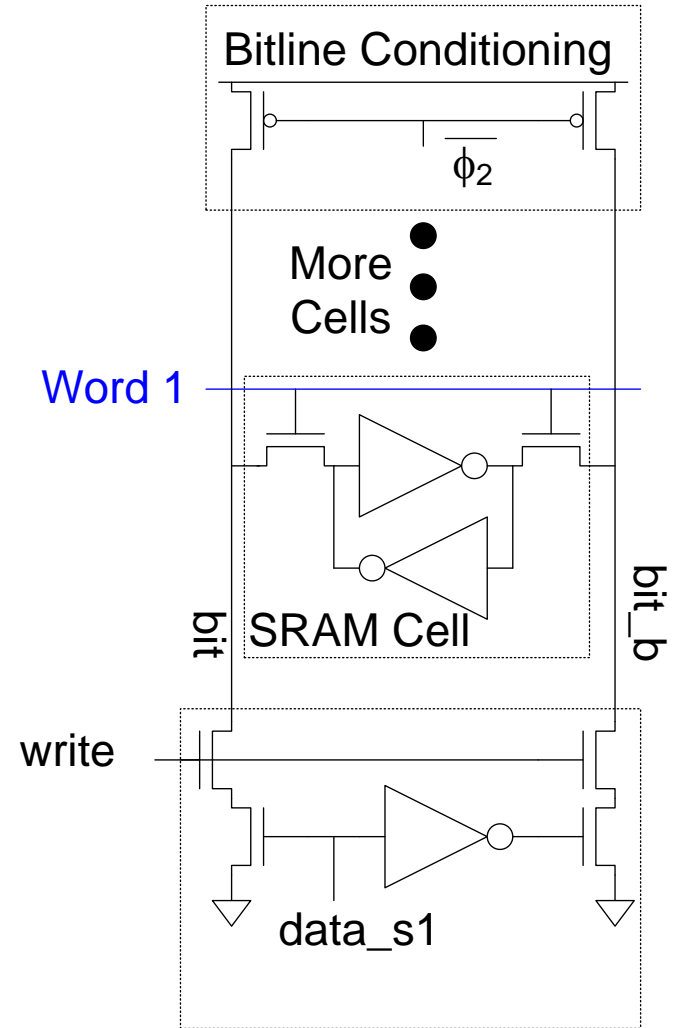


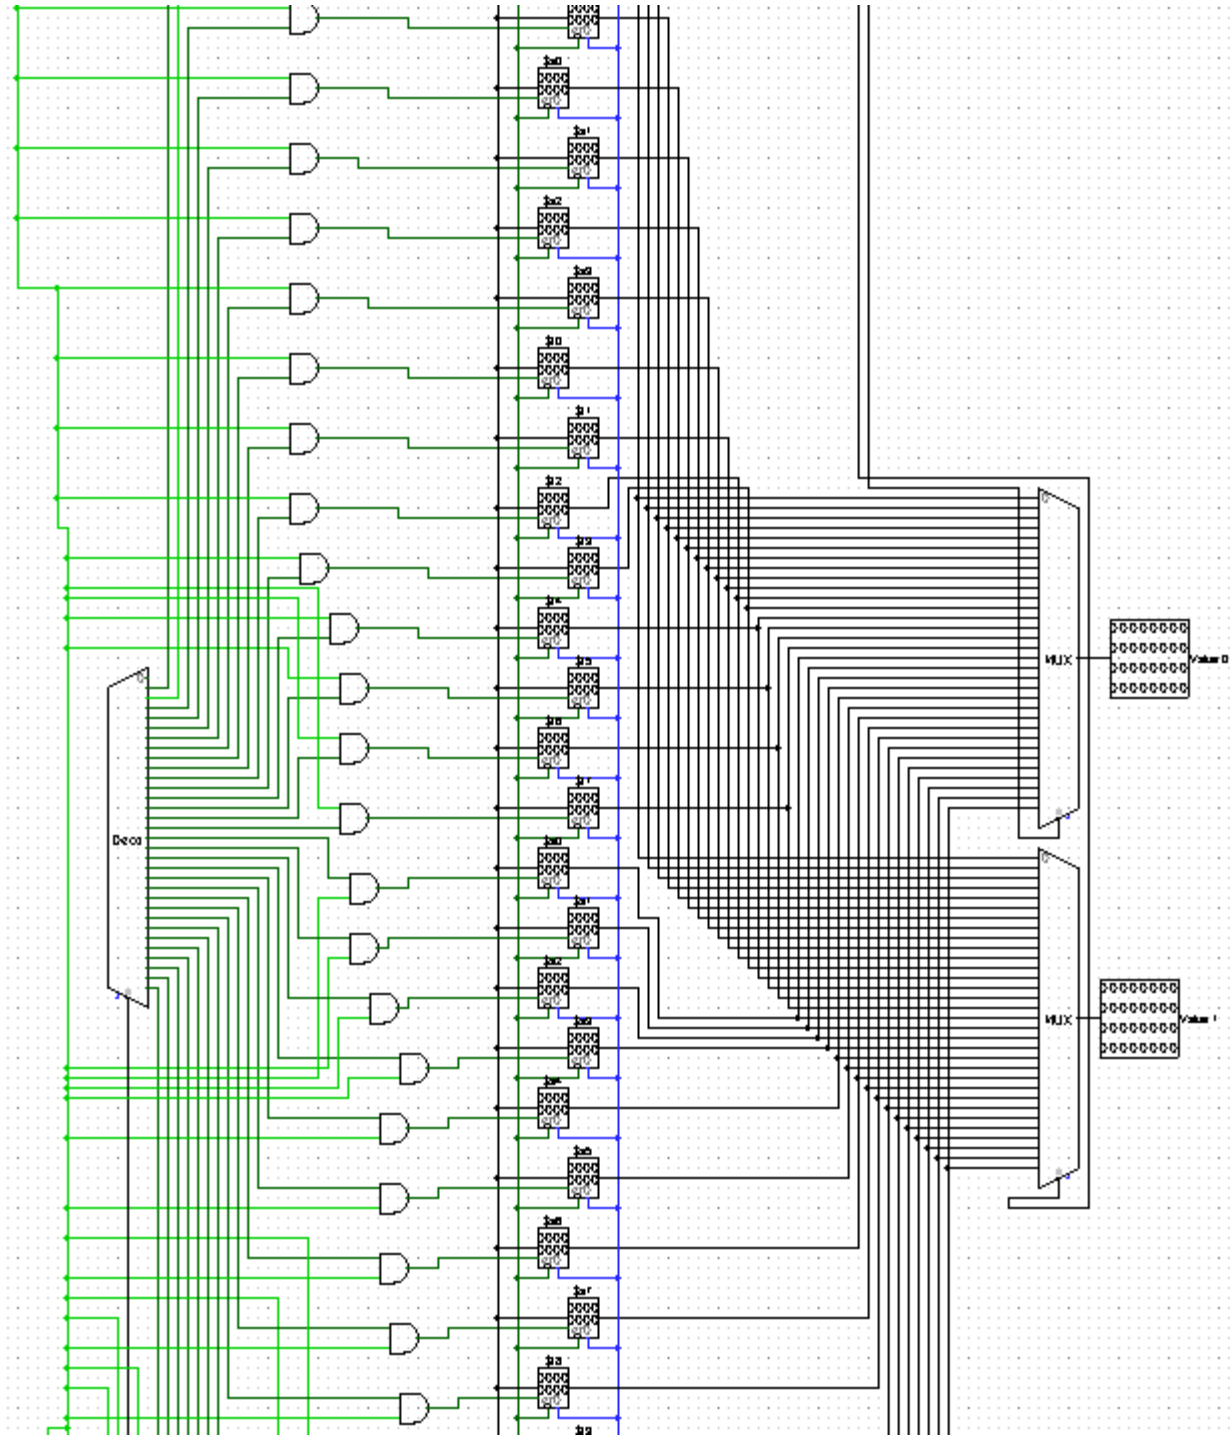
# SRAM Column Example

Read




Write






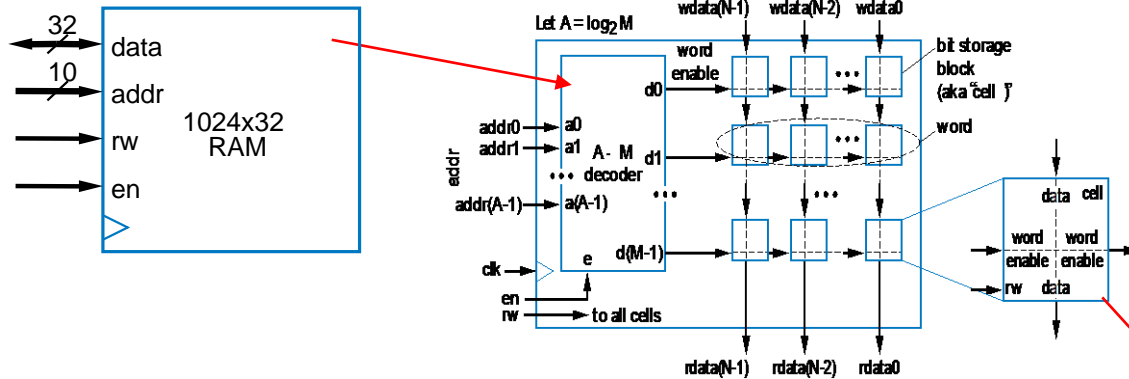
# SRAM 4x4 array



# SRAM 4x4 array

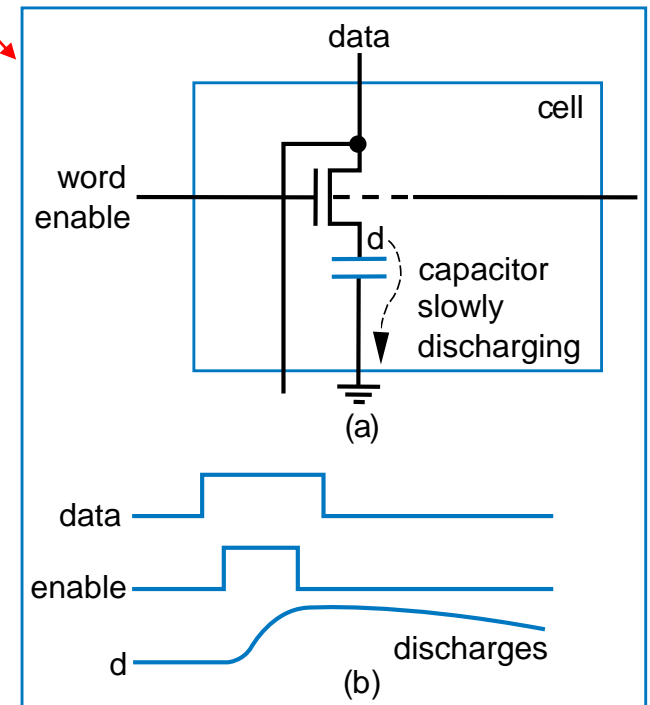


# Dynamic RAM (DRAM)



DRAM cell

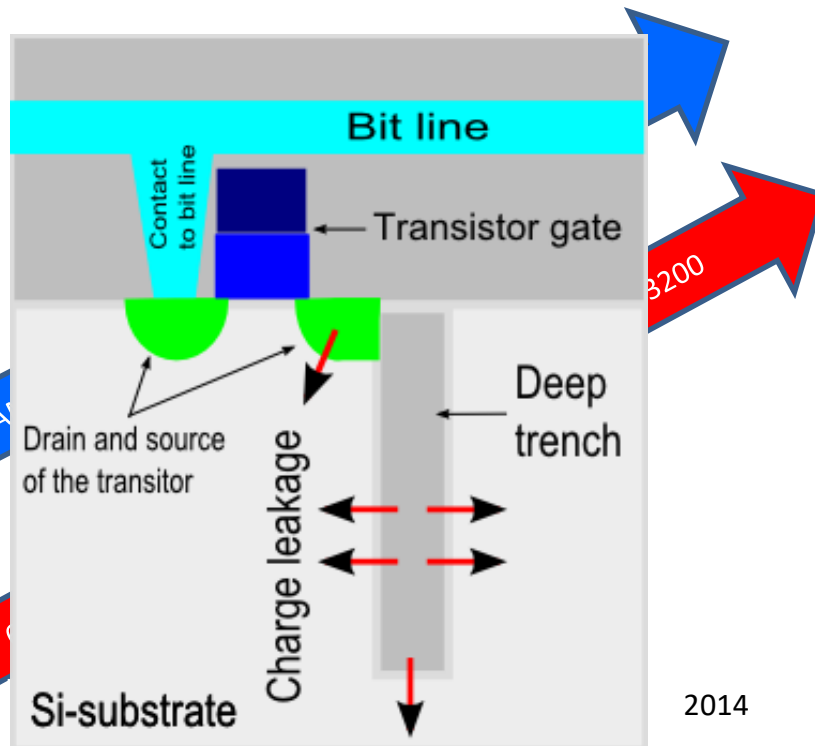
- “Dynamic” RAM cell
  - 1 transistor (rather than 6)
  - Relies on *large* capacitor to store bit
  - Write: Transistor conducts, data voltage level gets stored on top plate of capacitor
  - Read: Just look at value of  $d$
  - Problem: Capacitor discharges over time
    - Must “refresh” regularly, by reading  $d$  and then writing it right back



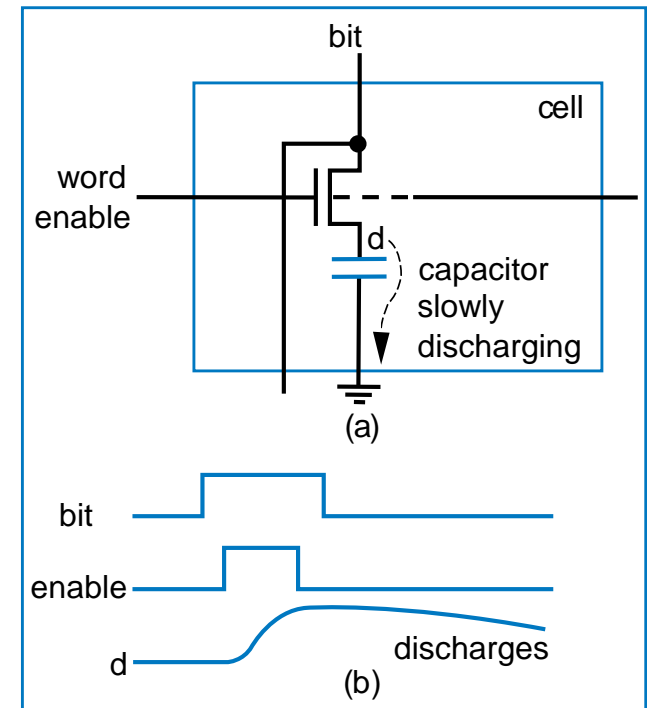


# Dynamic RAM (DRAM)

- “Dynamic” RAM cell
  - 1 transistor (rather than 6)
  - Relies on capacitor to store bit





DRAM cell




2014

# Dynamic RAM 1-Transistor Cell: Layout




Trench Cell



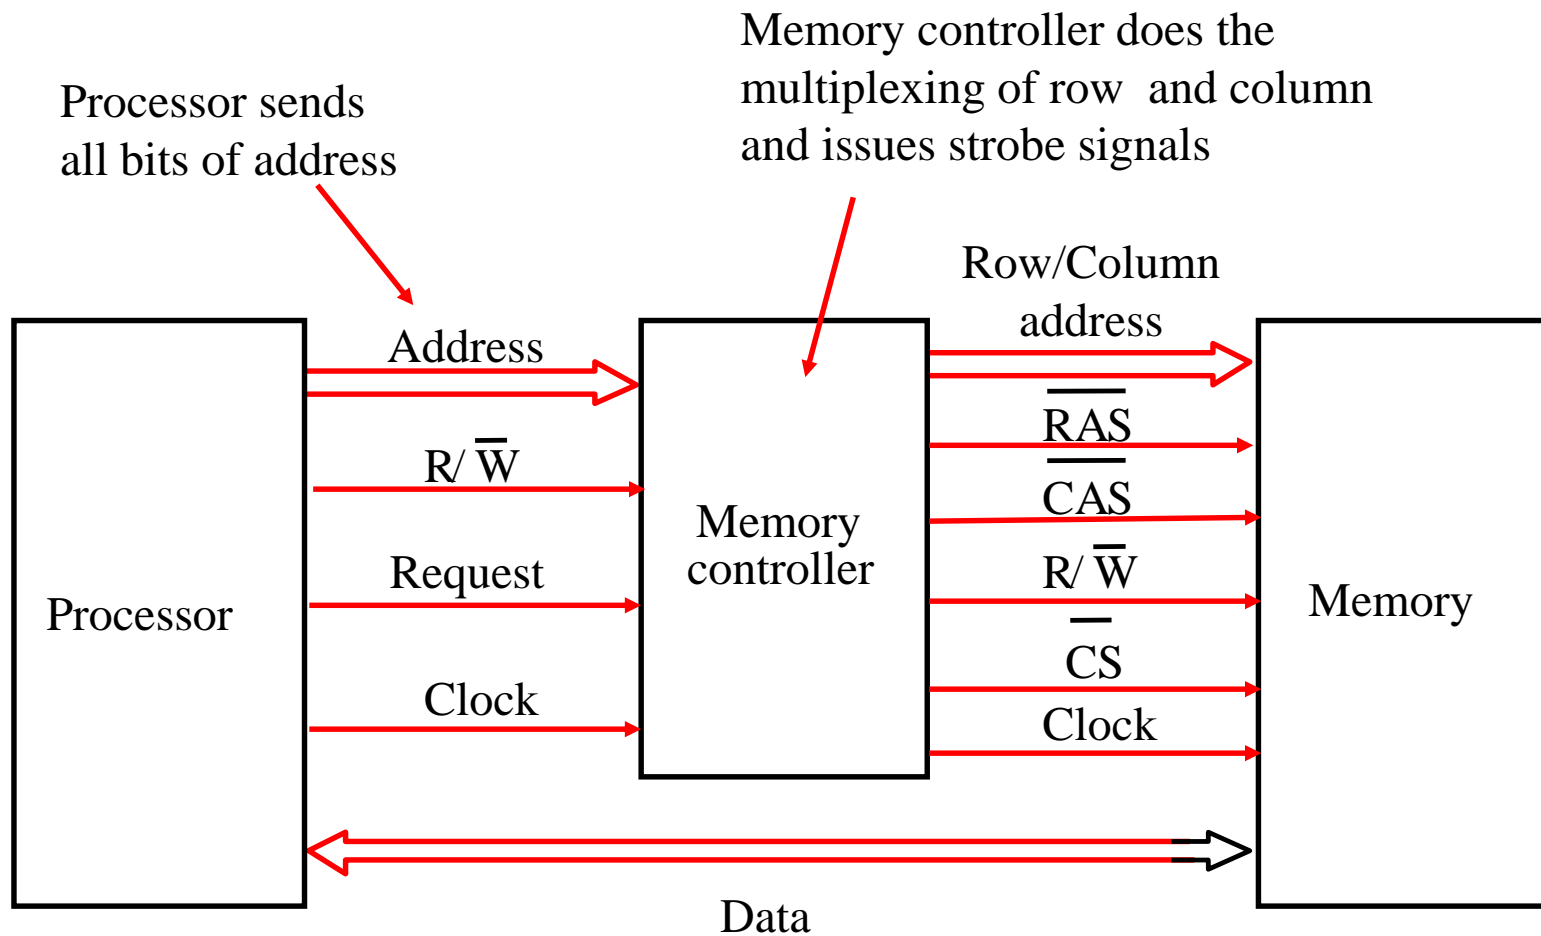
Stacked-capacitor Cell

# Comparing Memory

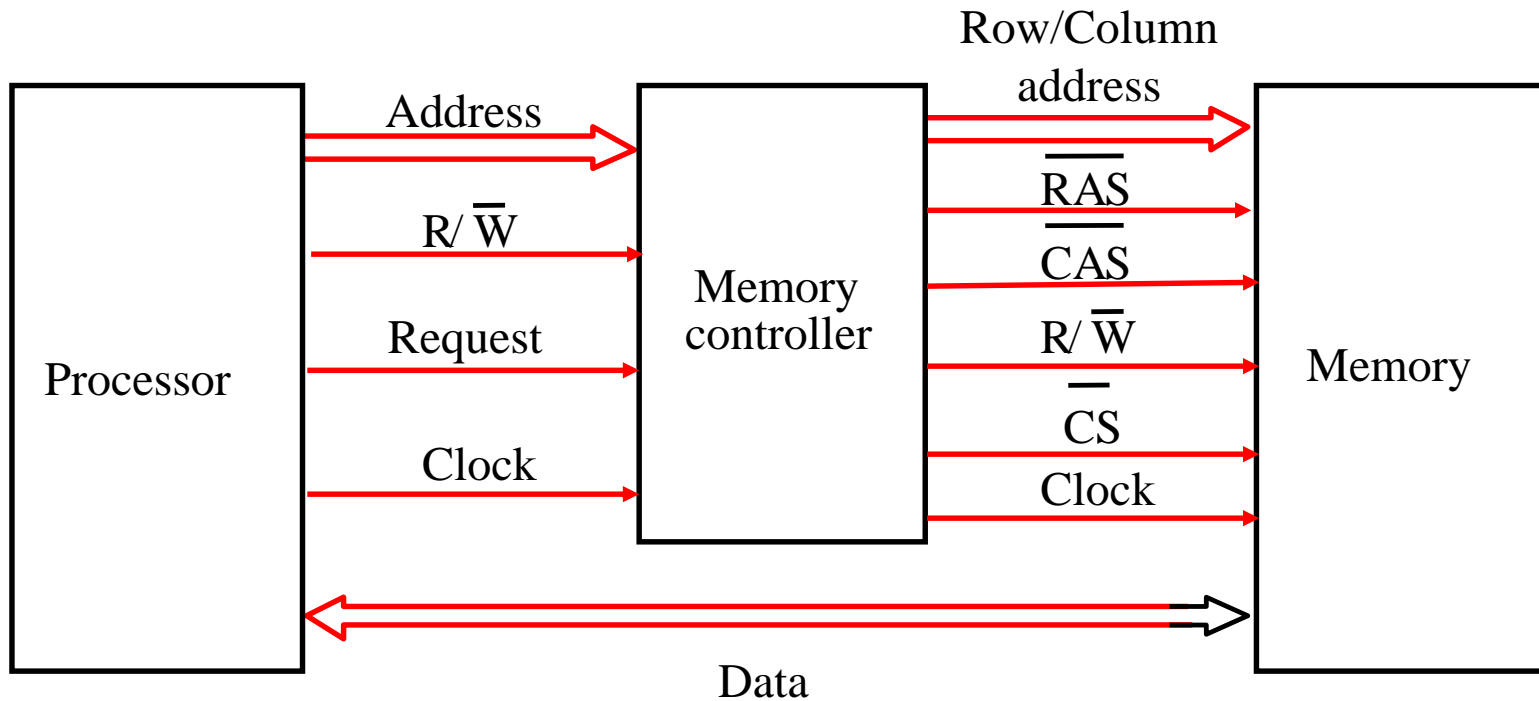
- Register file
  - Fastest
  - But biggest size
- SRAM
  - Fast
  - More compact than register file
- DRAM
  - Slowest (capacitor)
    - And refreshing takes time
  - But very compact (lower cost)
- Use register file for small items, SRAM for large items, and DRAM for huge items
  - Note: DRAM's big capacitor requires a special chip design process, so DRAM is often a separate chip



Size comparison for same number of bits (not to scale)



Use of a memory controller.



Memory controller provides the refresh control if not done on the chip


Refreshing typically once every 64 ms. At a cost of .2ms

Less than .4% overhead

Use of a memory controller.


# CAM

- CAM vs. RAM



# CAM


- Binary CAM Cell
  - ML pre-charged to  $V_{DD}$
  - Match: ML remains at  $V_{DD}$
  - Mismatch: ML discharges





# Read-Only Memory

- A block diagram of a ROM is shown below. It consists of  $k$  address inputs and  $n$  data outputs.
- The number of words in a ROM is determined from the fact that  $k$  address input lines are needed to specify  $2^k$  words.




ROM Block diagram  $2^k \times n$  Module

# Read-Only Memory Cells


Bit line (BL) is resistively clamped to the ground, so its default value is 0

Diode disadvantage – no electrical isolation between bit and word lines

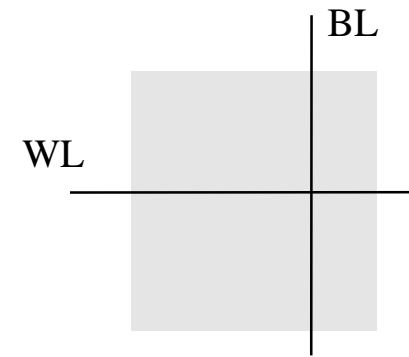
BL is resistively clamped to VDD, so its default value is 1






Diode ROM



MOS ROM 1




MOS ROM 2




# Nonvolatile Memory

ROM  
PROM  
EPROM




Read-only memory organization, with the fixed contents shown on the right.

# MOS NOR ROM




# MOS NAND ROM



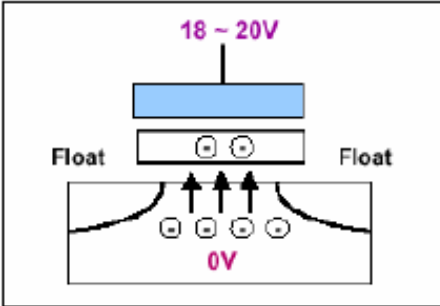
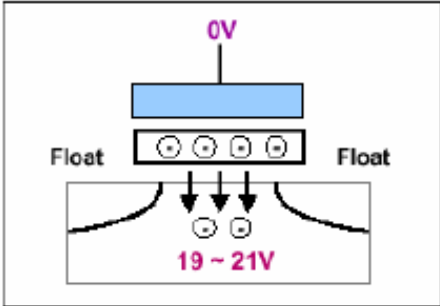
All word lines high by default with exception of selected row

# Flash Memory



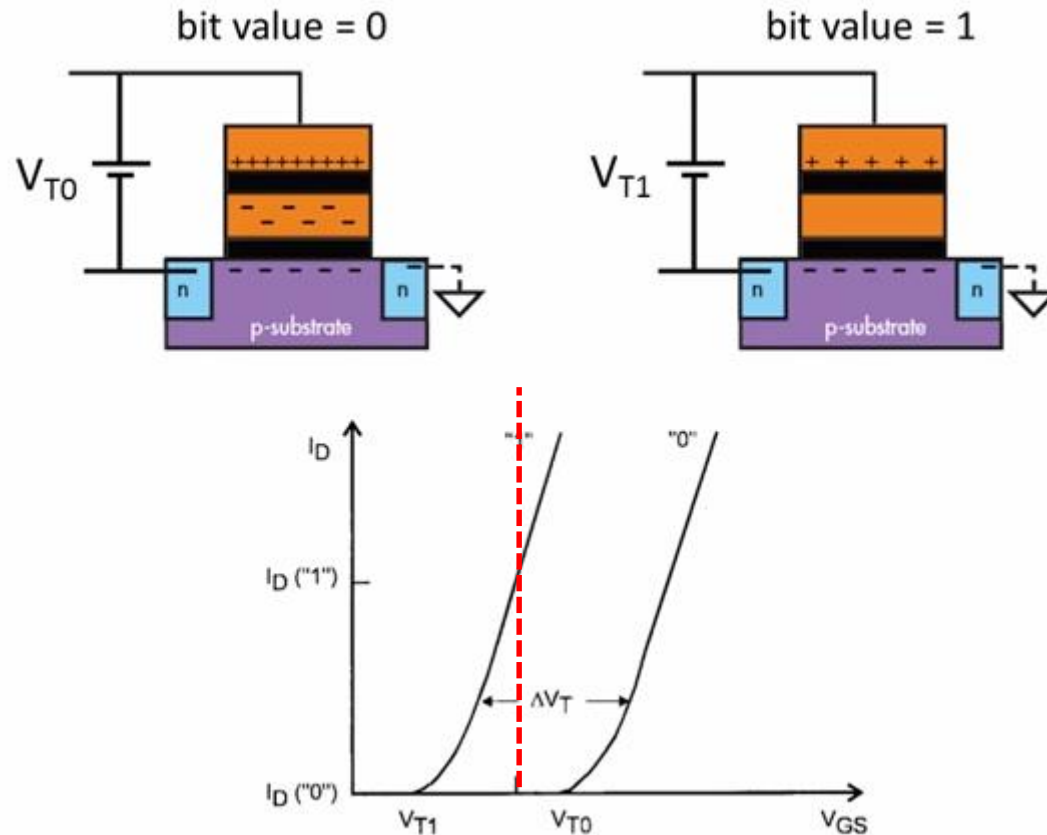
# Flash memory

- Flash memory
  - Non volatile read only memory (ROM)
  - Erase Electrically or UV (EPROM)
  - Uses F-N tunneling for program & erase
  - Reads like DRAM (~ns)
  - Writes like DISK (~ms).

Program	Erase
	
<ul style="list-style-type: none"><li>● Use F-N Tunneling</li><li>● Channel Inversion</li></ul>	<ul style="list-style-type: none"><li>● Use F-N Tunneling</li><li>● Channel Accumulation</li></ul>

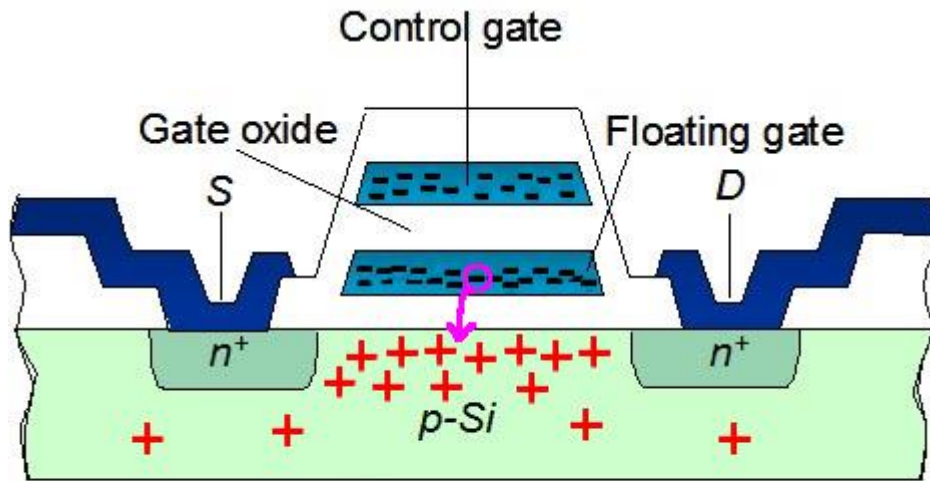


# Reading Memory State



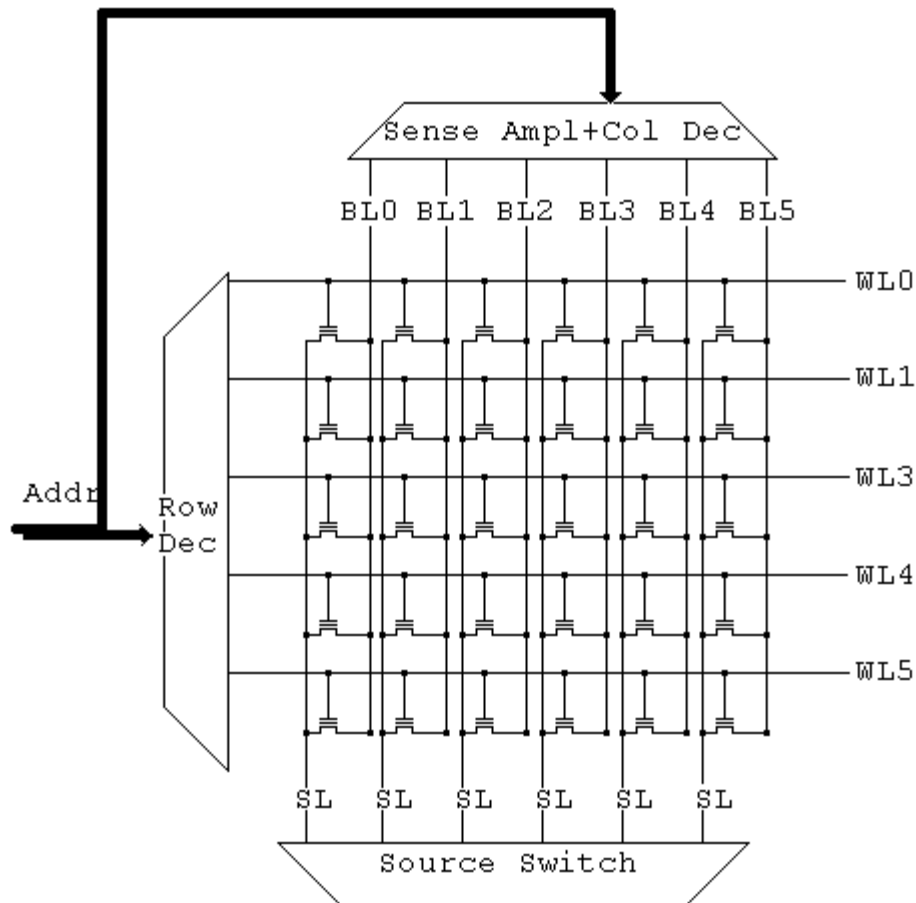
Change in Threshold Voltage due to **Screening Effect** of Floating Gate  
Read mode: Apply intermediate voltage, check whether current is flowing or not

# Writing Memory State



Control gate voltage determines whether electrons are injected to, or push/pulled out of floating gate.

# NOR Array




Reading:



Assert a single word line.  
The source lines are asserted and the read of the bitline gives the contents of the cell.

# Multi-Levels

- By using reference cells set at given levels and comparing them to the value from the bitline, we can determine the value stored.




# Review of memory technologies




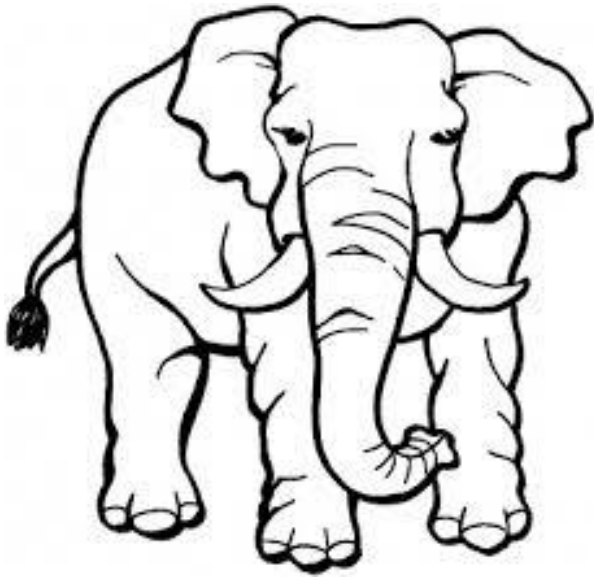
Memory type	SRAM	DRAM	Flash
Speed	Very fast	Slow	Slow
Density	Low	High	Very high
Power	High	Low	Very low
Refresh	No	Yes	No
Mechanism	Bi-stable latch	Capacitor	FN tunneling

# Memory Hierarchy

# Actual Memory Systems



+



Type1: Fast, Small

Can we achieve?: Fast, Large  
Type2: Slow, Large

(Cache Principle)